



2110625 - Data Science Architecture

Document Store

Asst.Prof. Natawut Nupairoj, Ph.D.
Department of Computer Engineering
Chulalongkorn University
natawut.n@chula.ac.th

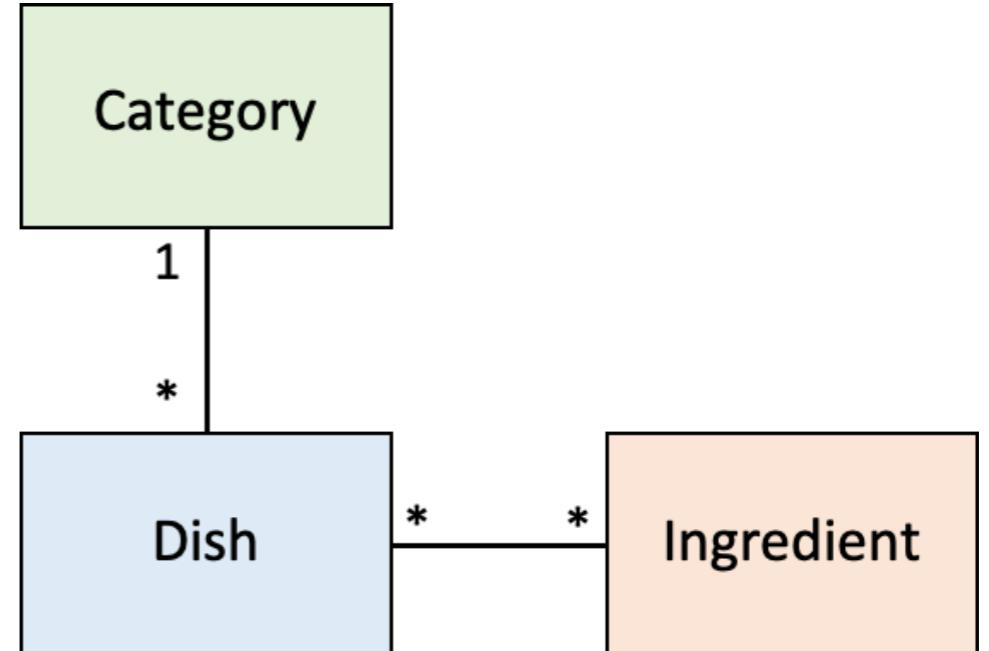
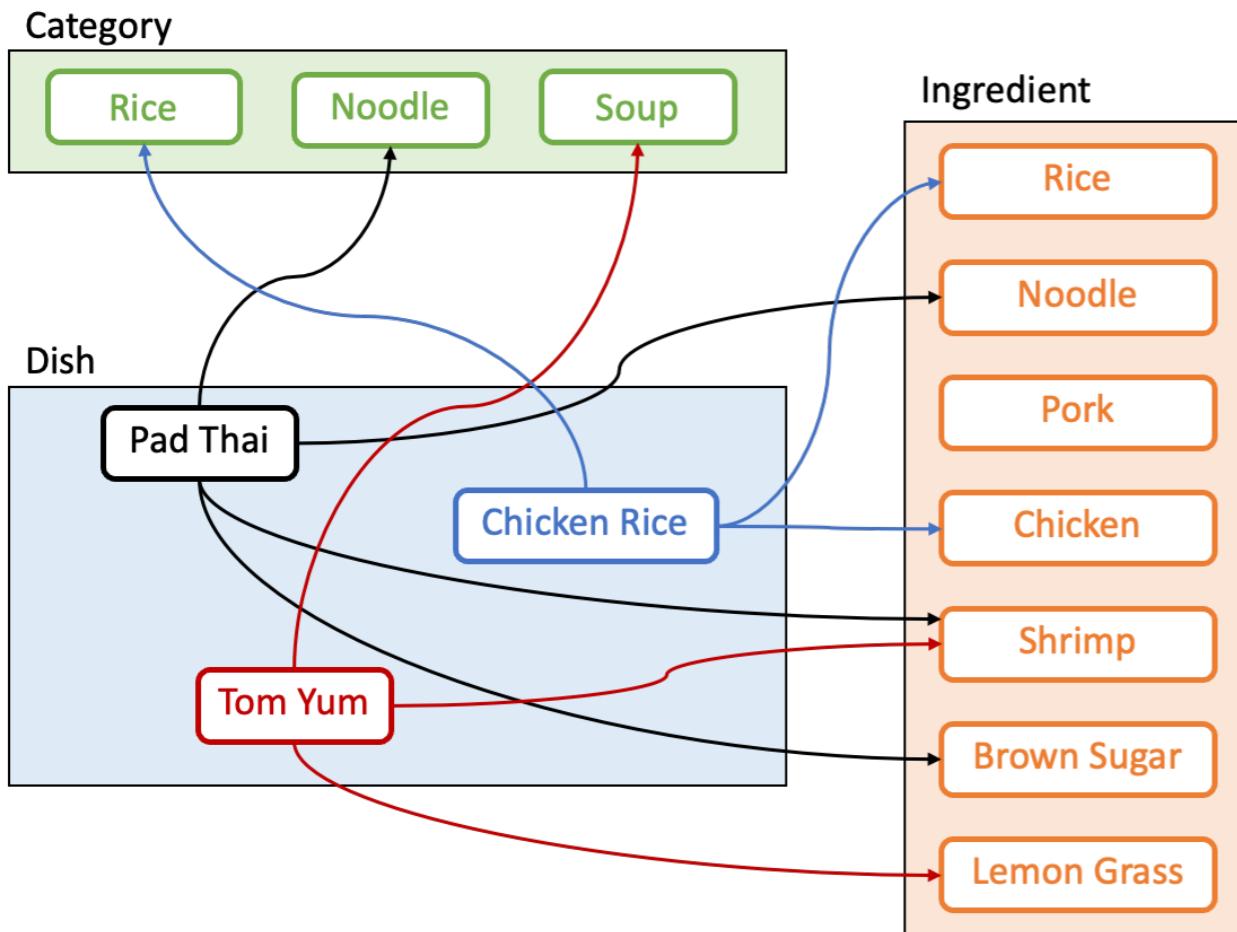
Document Store

ទីតាំង transaction base

- Data storage system designed for storing, retrieving and managing document (aka. semi-structured data)
- Data is stored as a document or item without any link or relation to other table
- Document store usually runs on a cluster of multiple (10s-1000s) storage nodes to allow horizontal scalability
 - Shards or partitions are fundamental mechanism for both scaling and replicated capabilities

កំណែ កំរែនទេស្ថុល

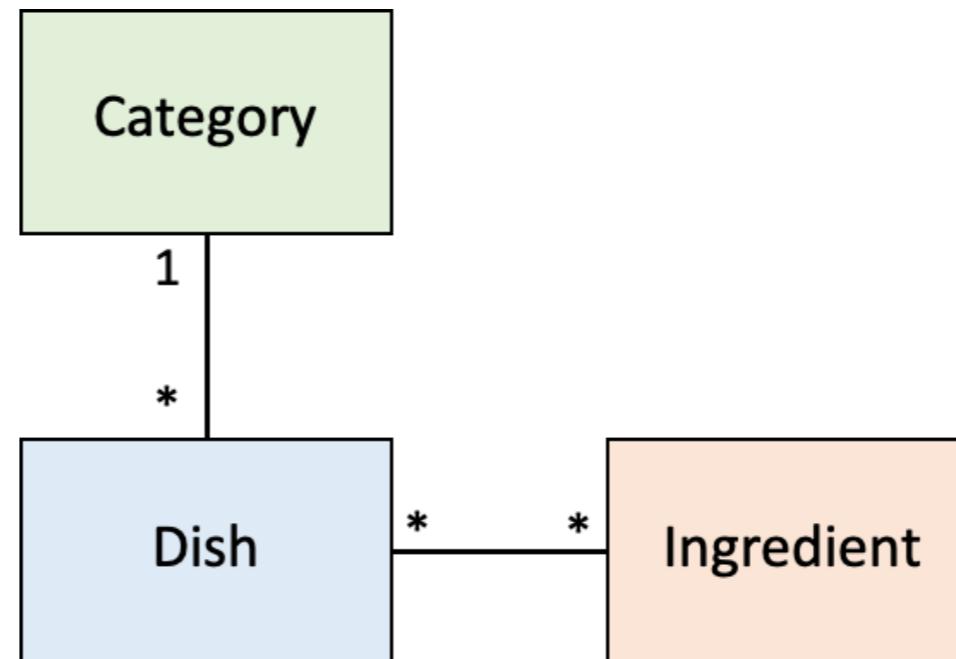
Recipe Database Revisited



Example RDBMS: Recipe

Category

id	category
1	Soup
2	Noodle
3	Rice



Dish

id	name	category	...
1203	Pad Thai	2	
1288	Chicken Rice	3	

DI_link

d_id	i_id	quantity	unit
1203	45	125	g
1203	48	150	g
1203	52	3	tbsp

Ingredient

id	item
45	noodle
46	rice
48	chicken
52	brown sugar

Example MongoDB: Recipe

Scale out មែនdependencies នៅក្នុងក្រុងការបង្កើត
នៅក្នុងក្រុងការបង្កើត doc ចាប់ពីលីអូន

```
{  
    _id: ObjectId("507f191e810c19729de860ea"),  
    name: "Pad Thai",  
    cuisine: "Thai",  
    category: "Noodle",  
    rating: 4.7,  
    votes: 9301,  
    ingredients: [  
        { item: "noodle", quantity: "125", unit: "g" },  
        { item: "chicken", quantity: "150", unit: "g" },  
        { item: "brown sugar", quantity: "3", unit: "tbsp" }  
        ...  
    ]  
}
```



Document Store
MongoDB

MongoDB

- Document-Oriented NoSQL database
- Based on CP concept *CP* ເນັ້ນ *Consistency* ແລະ *scale out* ໄດ້
- BSON store (binary-format JSON)
- Databases – Collections - Documents
- Schemaless - documents in the same collections can have different structures *ໄປຈາກສິນທີຖຸກ record ສໍາ field ເກມອນກົມ (ຍາກວັນ required ແລະ alpha id)*
- Query using javascript functions
- Support load-balancing, replication, scale-out, etc.

RDBMS vs. MongoDB

RDBMS

Database

Table

Row / Tuple

Column

Table Join *ఫ్లోవర్స్ ను scale out*

Primary Key

MongoDB

Database

Collection

Document

Field

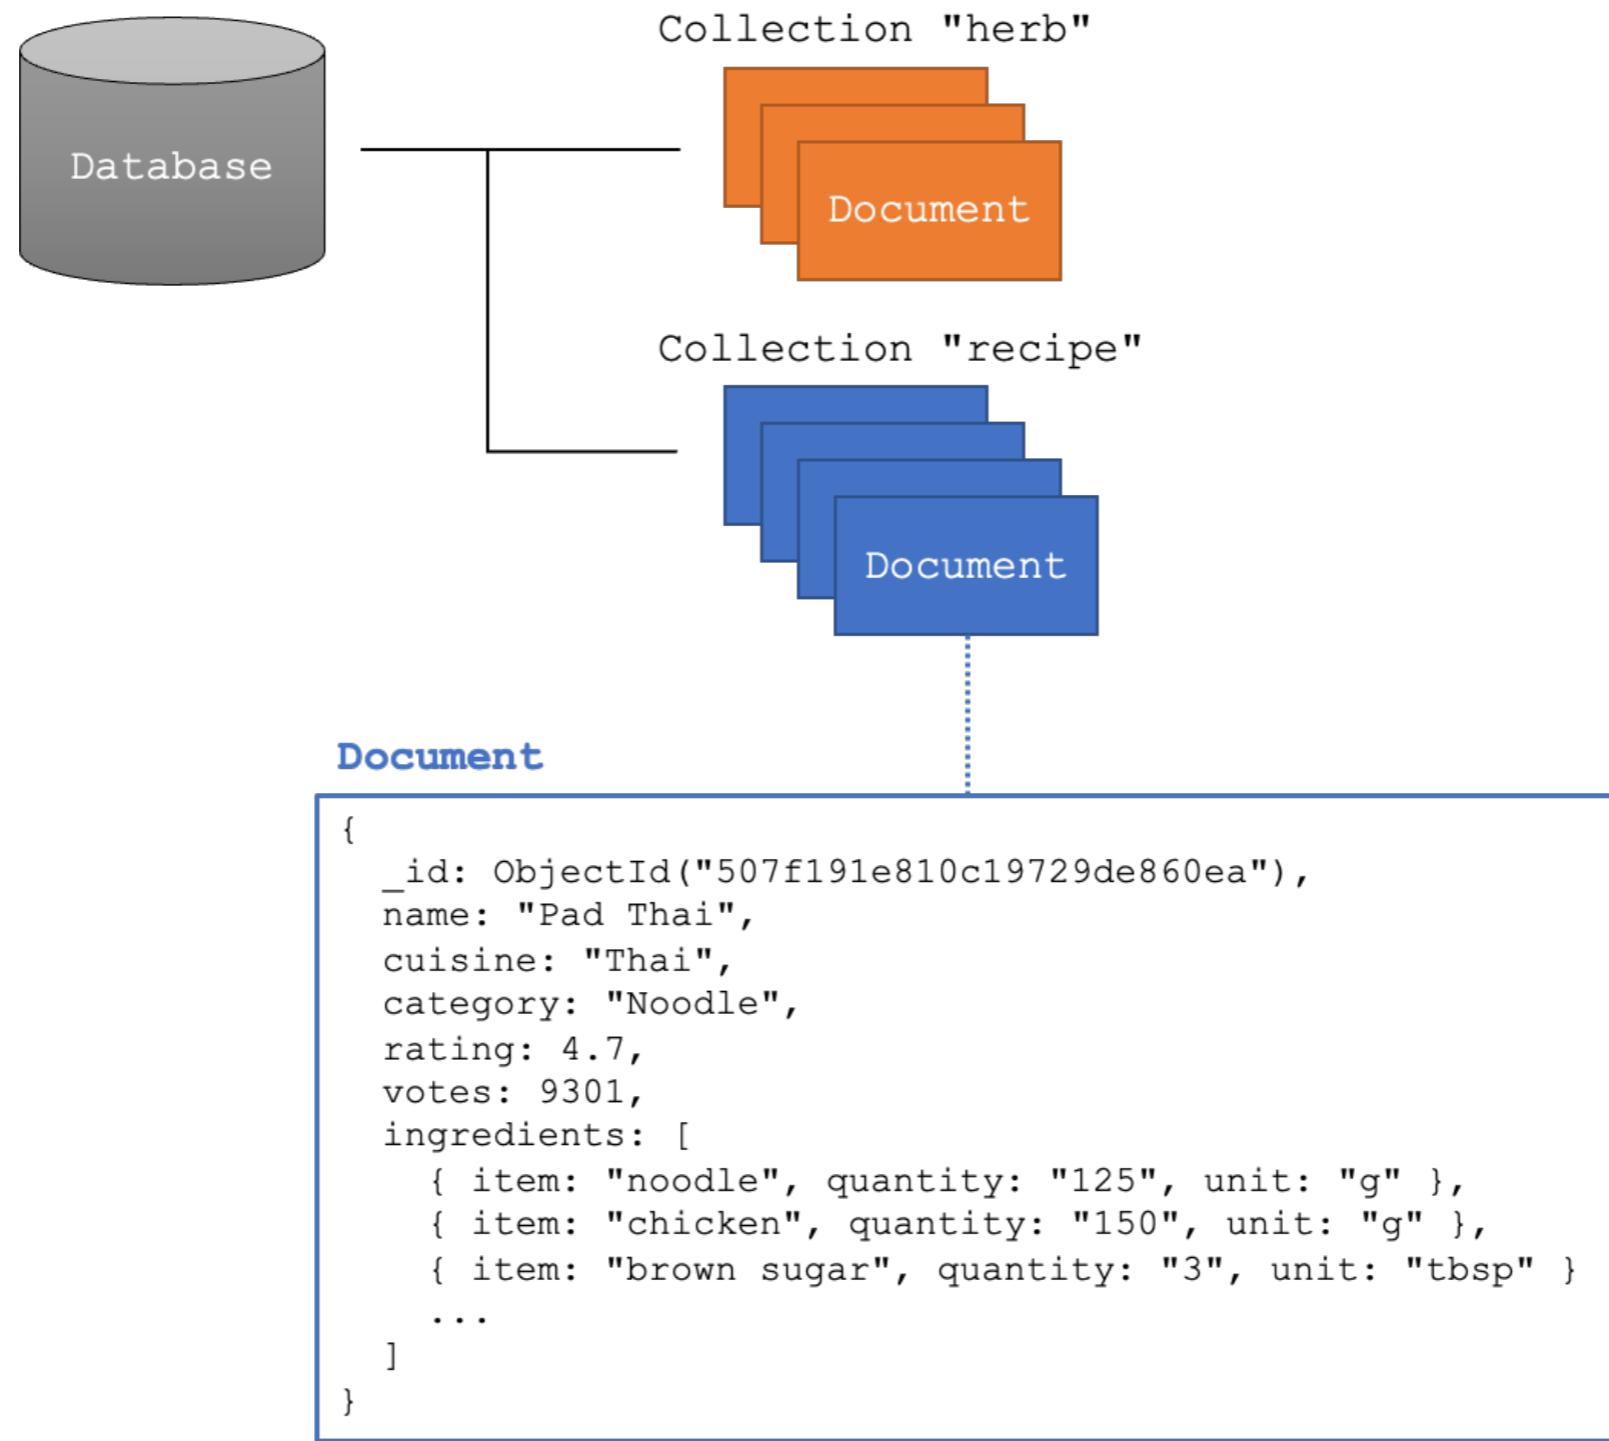
Embedded Document

Primary Key
(Default key “_id” provided by mongodb itself)

Example Document: Recipe

```
{  
  _id: ObjectId("507f191e810c19729de860ea"),  
  name: "Pad Thai",  
  cuisine: "Thai",  
  category: "Noodle",  
  rating: 4.7,  
  votes: 9301,  
  ingredients: [  
    { item: "noodle", quantity: "125", unit: "g" },  
    { item: "chicken", quantity: "150", unit: "g" },  
    { item: "brown sugar", quantity: "3", unit: "tbsp" }  
    ...  
  ]  
}
```

Basic Terminologies



MongoDB CRUD Operations - Insert

```
1 use("restaurants")                                     collection "recipe"
2
3 db.recipe.insertOne(                                 ←
4 {                                                 ←
5   name: "Pad Thai",                            field: value
6   cuisine: "Thai",                            field: value
7   rating: 4.7,                                field: value
8   votes: 9301,
9   ingredients: [                                ←
10    { item: "noodle", quantity: "125", unit: "g" },
11    { item: "chicken", quantity: "150", unit: "g" },
12    { item: "brown sugar", quantity: "3", unit: "tbsp" }
13  ]
14 }
15 )
```

document

MongoDB CRUD Operations - Query and Others

```
db.recipe.find(  
  { cuisine: "Thai", rating: { $gt: 4.4 } }, ← query criteria  
  { name: 1, rating: 1, category: 1 } ← projection  
).limit(2)   ← cursor modifier  
             ↙ select name, rating, category
```

- Other CRUD operations

```
db.collection.updateOne()  
db.collection.updateMany()  
db.collection.replaceOne()  
db.collection.deleteOne()  
db.collection.deleteMany()
```

- See <https://www.mongodb.com/docs/manual/crud/> for more information

More on Query Operation

Collection

```
db.recipe.find( {rating: {$gt: 4}} ).sort( {rating: -1} )
```

{ rating: 4.5, ... }
{ rating: 4.1, ... }
{ rating: 4.9, ... }
{ rating: 2.7, ... }
{ rating: 4.3, ... }
{ rating: 3.8, ... }

recipe

Query Criteria

Query
Criteria

{ rating: 4.5, ... }
{ rating: 4.1, ... }
{ rating: 4.9, ... }
{ rating: 4.3, ... }

Modifier

Modifier

{ rating: 4.9, ... }
{ rating: 4.5, ... }
{ rating: 4.3, ... }
{ rating: 4.1, ... }

Results

Aggregation Operations (sum, avg) խցկնիղն կամ aggregate հելութիւն

Collection

```
db.recipe.aggregate( [  
  $match stage → { $match: { rating: { $gt: 3.0 } } },  
  $group stage → { $group: { _id: "$category", total: { $sum: "$votes" } } }  
]
```

{	region: "Noodle",	
	rating: 4.5,	
	votes: 508	
}		
{	region: "Seafood",	
	rating: 4.1,	
	votes: 1832	
}		
{	region: "Dessert",	
	rating: 2.7,	
	votes: 295	
}		
{	region: "Seafood",	
	rating: 4.3,	
	votes: 14	
}		
{	region: "Noodle",	
	rating: 3.8,	
	votes: 931	
}		

rating > 3
\$match

{	region: "Noodle",	
	rating: 4.5,	
	votes: 508	
}		
{	region: "Seafood",	
	rating: 4.1,	
	votes: 1832	
}		
{	region: "Seafood",	
	rating: 4.3,	
	votes: 14	
}		
{	region: "Noodle",	
	rating: 3.8,	
	votes: 931	
}		

sum (vote)
\$group

{	_id: "Noodle",	
	total: 1439	
}		
{	_id: "Seafood",	
	total: 1846	
}		

Փոխարին SQL հայելեակը ըստ շաբաթական գործության

Playing with MongoDB - Client

- For client-side, you have a couple options
 - Using MongoDB (JavaScript) with VS Code, install MongoDB for VS Code from the VS Code Marketplace
 - Using MongoDB with python, install pymongo
- pip install pymongo

Playing with MongoDB - Server (container)

- Use official mongodb image at
https://hub.docker.com/_/mongo

```
docker run -d --name mongo --rm -p 27017:27017 mongo
```

- You can also clone https://github.com/natawutn/datasci_architecture repo for docker compose and example files
 - To run single server (with mongo express), use docker-compose.yml in compose/single-server folder

```
12 // Select the database to use.  
13 use('mongodbVSCodePlaygroundDB');  
14  
15 // Insert a few documents into the sales collection.  
16 db.getCollection('sales').insertMany([  
17   { 'item': 'abc', 'price': 10, 'quantity': 2, 'date': new Date('2014-03-01T08:00:00Z') },  
18   { 'item': 'jkl', 'price': 20, 'quantity': 1, 'date': new Date('2014-03-01T09:00:00Z') },  
19   { 'item': 'xyz', 'price': 5, 'quantity': 10, 'date': new Date('2014-03-15T09:00:00Z') },  
20   { 'item': 'xyz', 'price': 5, 'quantity': 20, 'date': new Date('2014-04-04T11:21:39.736Z') },  
21   { 'item': 'abc', 'price': 10, 'quantity': 10, 'date': new Date('2014-04-04T21:23:13.331Z') },  
22   { 'item': 'def', 'price': 7.5, 'quantity': 5, 'date': new Date('2015-06-04T05:08:13Z') },  
23   { 'item': 'def', 'price': 7.5, 'quantity': 10, 'date': new Date('2015-09-10T08:43:00Z') },  
24   { 'item': 'abc', 'price': 10, 'quantity': 5, 'date': new Date('2016-02-06T20:20:13Z') },  
25 ]);  
26  
27 // Run a find command to view items sold on April 4th, 2014.  
28 const salesOnApril4th = db.getCollection('sales').find({  
29   date: { $gte: new Date('2014-04-04'), $lt: new Date('2014-04-05') }  
30 }).count(); >=   
31  
32 // Print a message to the output window.  
33 console.log(` ${salesOnApril4th} sales occurred in 2014. `);  
34  
35 // Here we run an aggregation and open a cursor to the results.  
36 // Use '.toArray()' to exhaust the cursor to return the whole result set.  
37 // You can use '.hasNext()/next()' to iterate through the cursor page by page.  
38 db.getCollection('sales').aggregate([  
39   // Find all of the sales that occurred in 2014.  
40   { $match: { date: { $gte: new Date('2014-01-01'), $lt: new Date('2015-01-01') } } },  
41   // Group the total sales for each product.  
42   { $group: { _id: '$item', totalSaleAmount: { $sum: { $multiply: [ '$price', '$quantity' ] } } } }  
43 ]);
```



Python and MongoDB Introduction Notebook

This notebook bases on a mongodb tutorial blog "Getting Started with Python and MongoDB" available at <https://www.mongodb.com/blog/post/getting-started-with-python-and-mongodb>

```
In [ ]: from pymongo import MongoClient
from bson.objectid import ObjectId
from random import randint
from pprint import pprint

# Connect to MongoDB, change the << MONGODB URL >> to reflect your own connection string
client = MongoClient()
# Connect to collection - Note: Change connection string as needed
db = client.restaurants
```

Basic operations: insert, query, and delete

```
In [ ]: doc = { 'name': 'Juicy Burger', 'rating': 4, 'cuisine': 'American' }
result = db.reviews.insert_one(doc)
oid = result.inserted_id
print(oid)
```

```
In [ ]: for doc in db.reviews.find():
    print(doc)
```

```
In [ ]: db.reviews.delete_one({'_id': ObjectId(oid)})
```

```
In [ ]: for doc in db.reviews.find():
    print(doc)
```

Generate sample data

```
In [ ]: #Step 2: Create sample data
names = ['Kitchen', 'Animal', 'State', 'Tastey', 'Big', 'City', 'Fish', 'Pizza', 'Goat', 'Salty', 'Sandwich', 'Lazy', 'Fun']
company_type = ['LLC', 'Inc', 'Company', 'Corporation']
company_cuisine = ['Pizza', 'Bar Food', 'Fast Food', 'Italian', 'Mexican', 'American', 'Sushi Bar', 'Vegetarian']
for x in range(1, 501):
    restaurant = {
        'name' : names[randint(0, (len(names)-1))] + ' ' + names[randint(0, (len(names)-1))] + ' ' + company_type[randint(0, (len(company_type)-1))],
        'rating' : randint(1, 5),
        'cuisine' : company_cuisine[randint(0, (len(company_cuisine)-1))]}
    #Step 3: Insert business object directly into MongoDB via insert_one
    result = db.reviews.insert_one(restaurant)
    #Step 4: Print to the console the ObjectId of the new document
    print('Created {0} of 500 as {1}'.format(x,result.inserted_id))
    #Step 5: Tell us that you are done
```

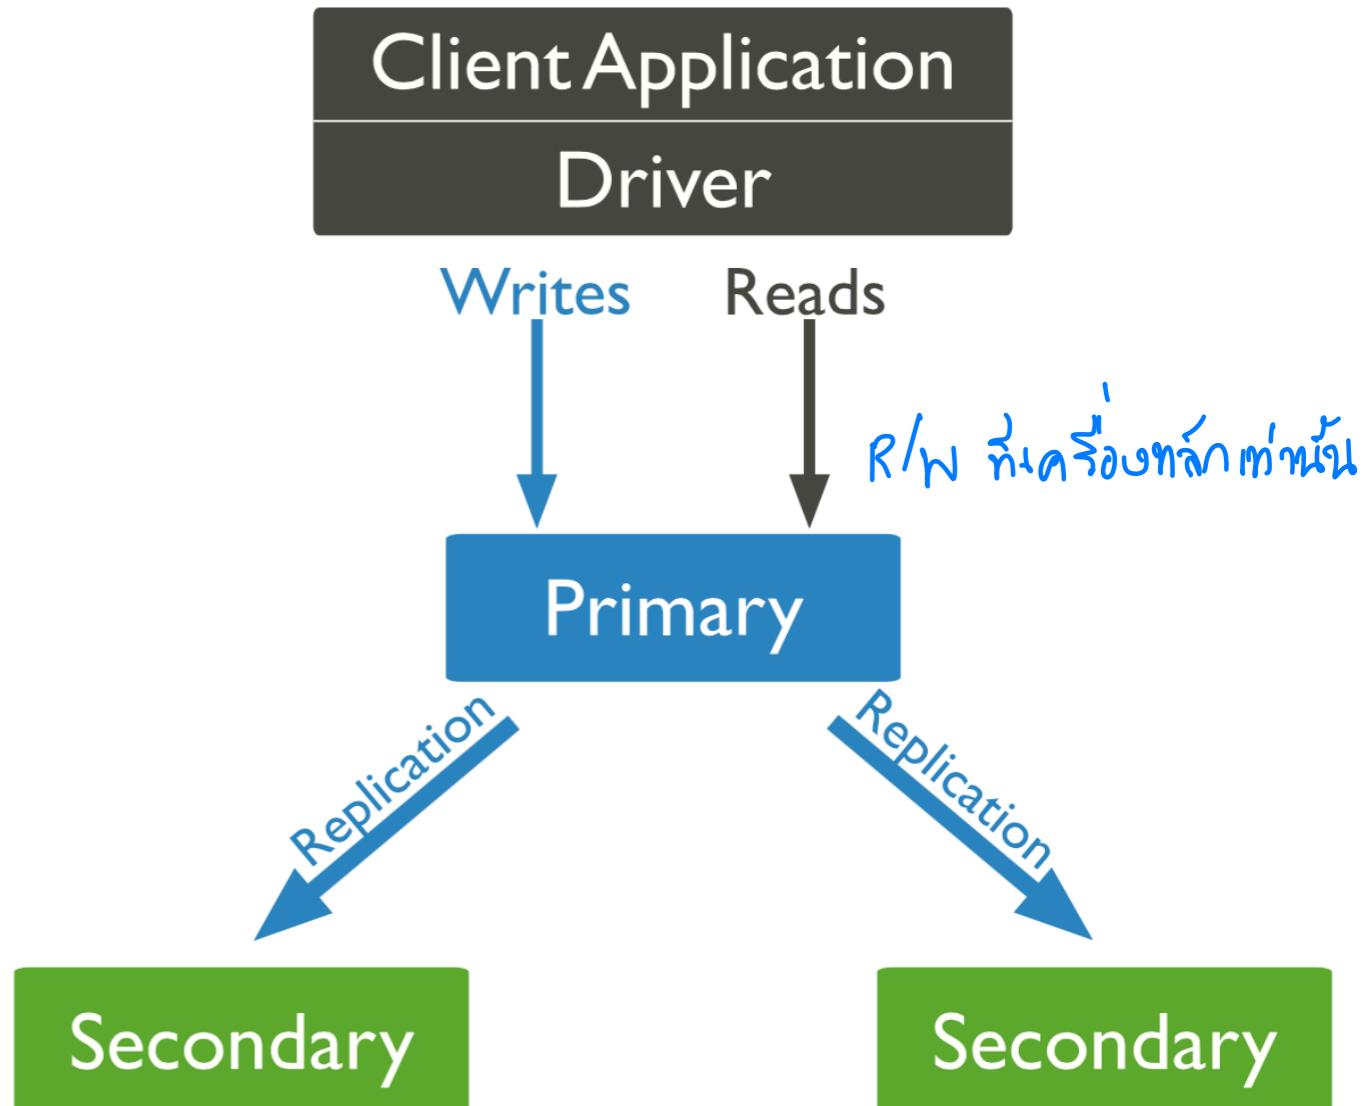
Highly Available with Replica Set

set នៃ server instance

- A group of mongod processes that maintain the same data set
- Provide redundancy and high availability
- Provide increased read capacity as clients can send read operations to different servers (in some cases)

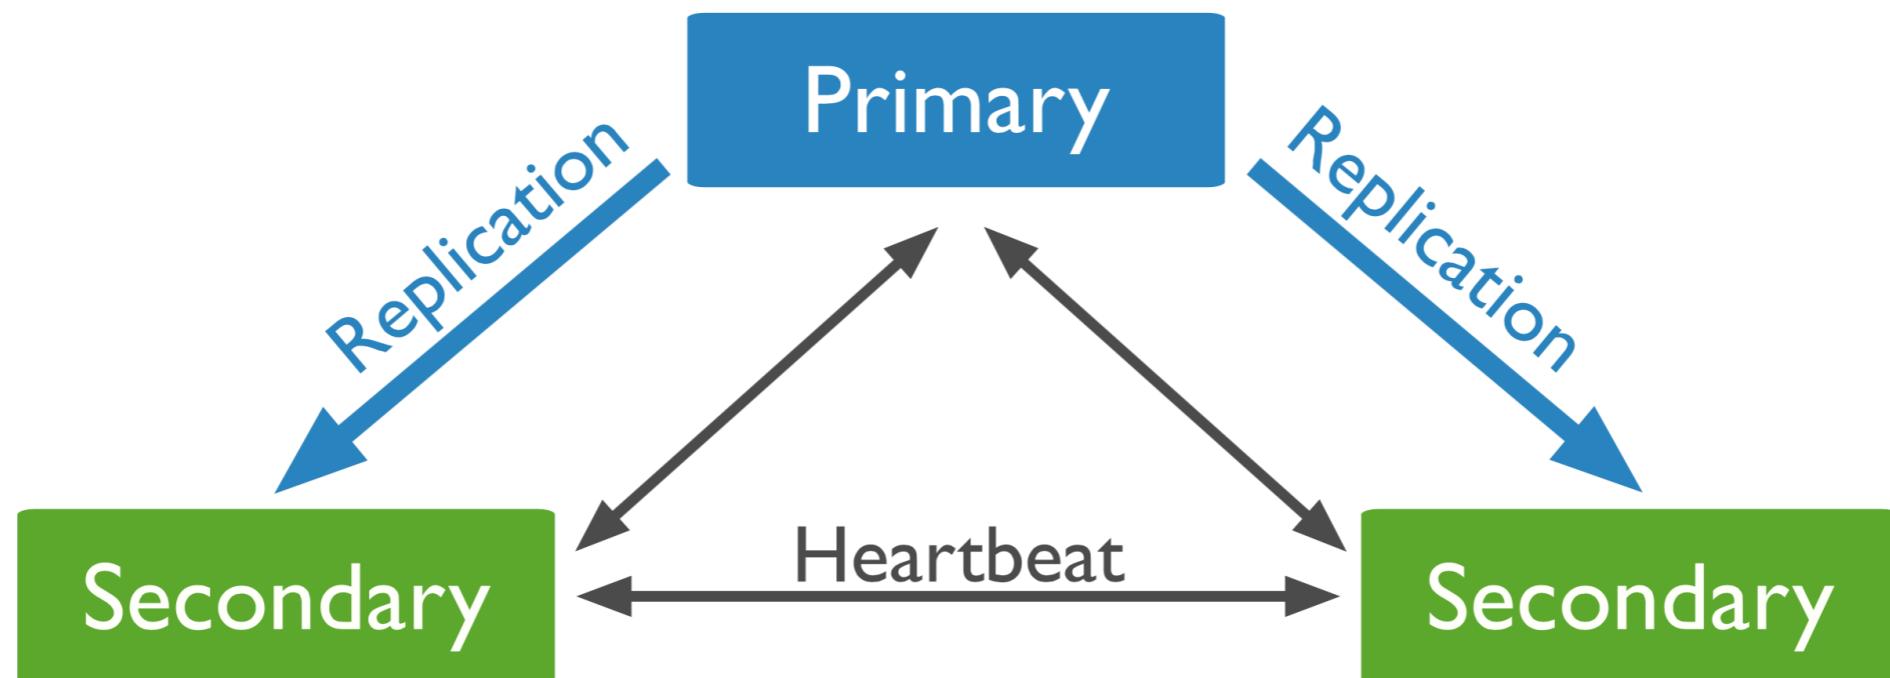
replica set = 3

គេ 1 គេ primary 2 គេ secondary

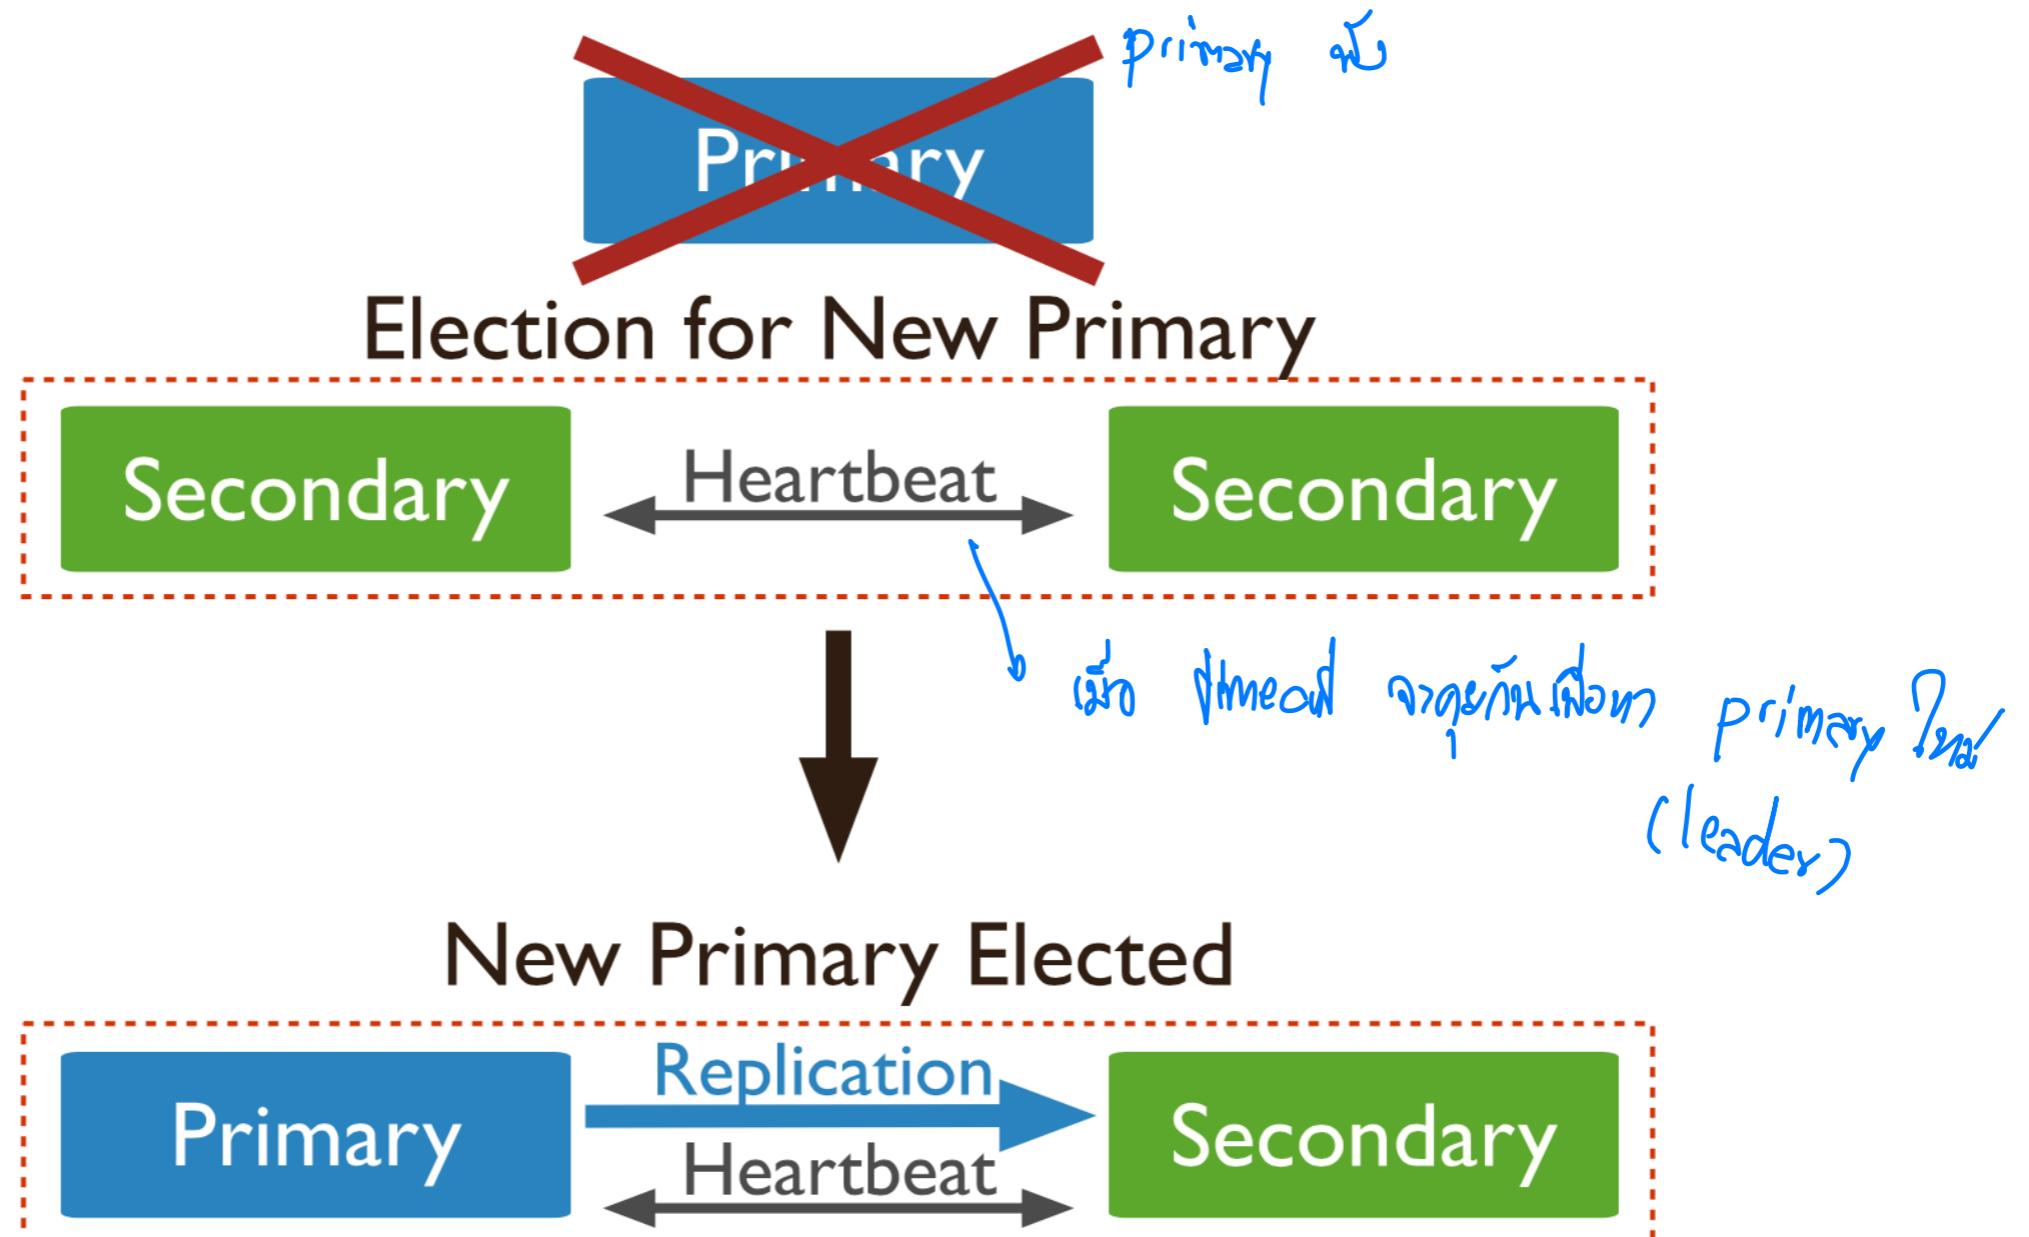


Source: <https://docs.mongodb.com/manual/replication/>

Replica Set in Actions



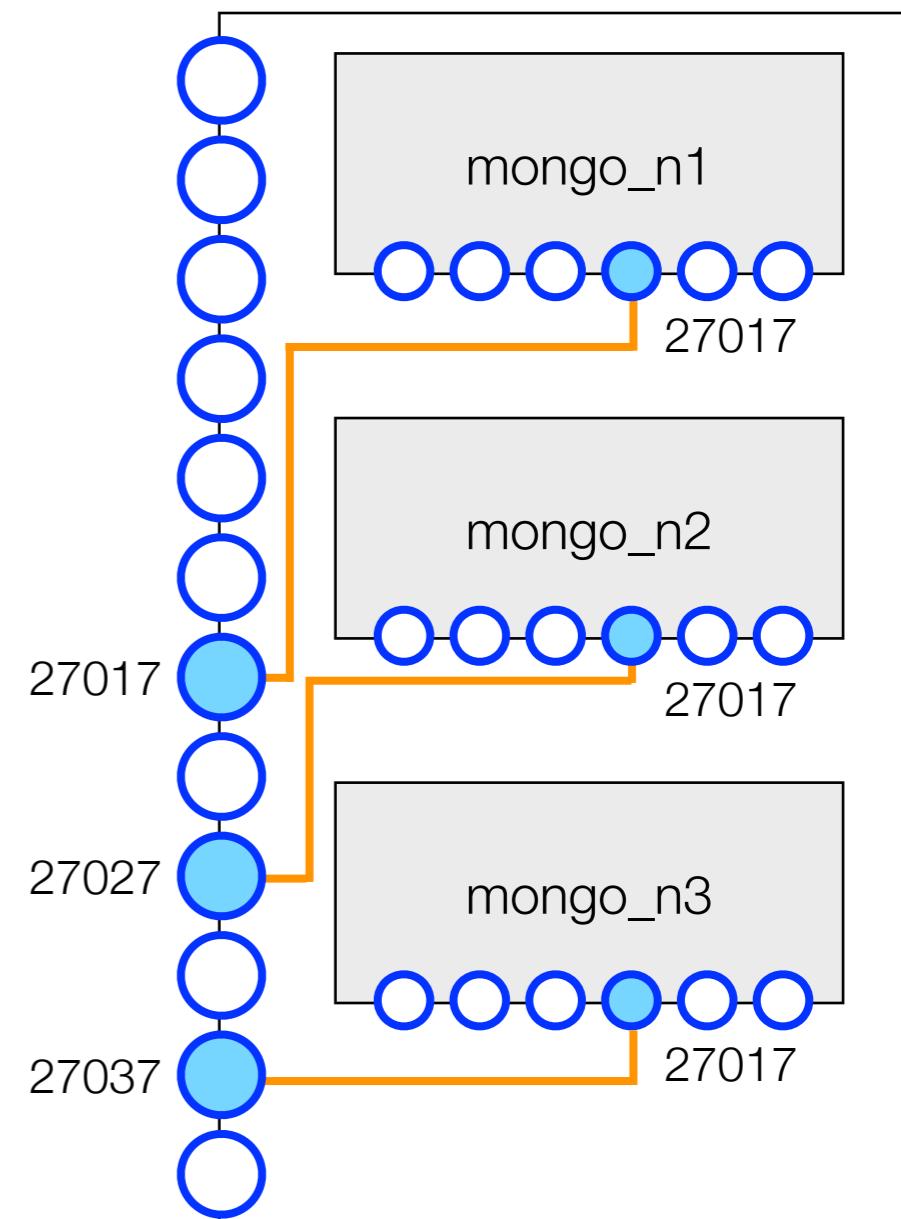
Replica Set in Actions



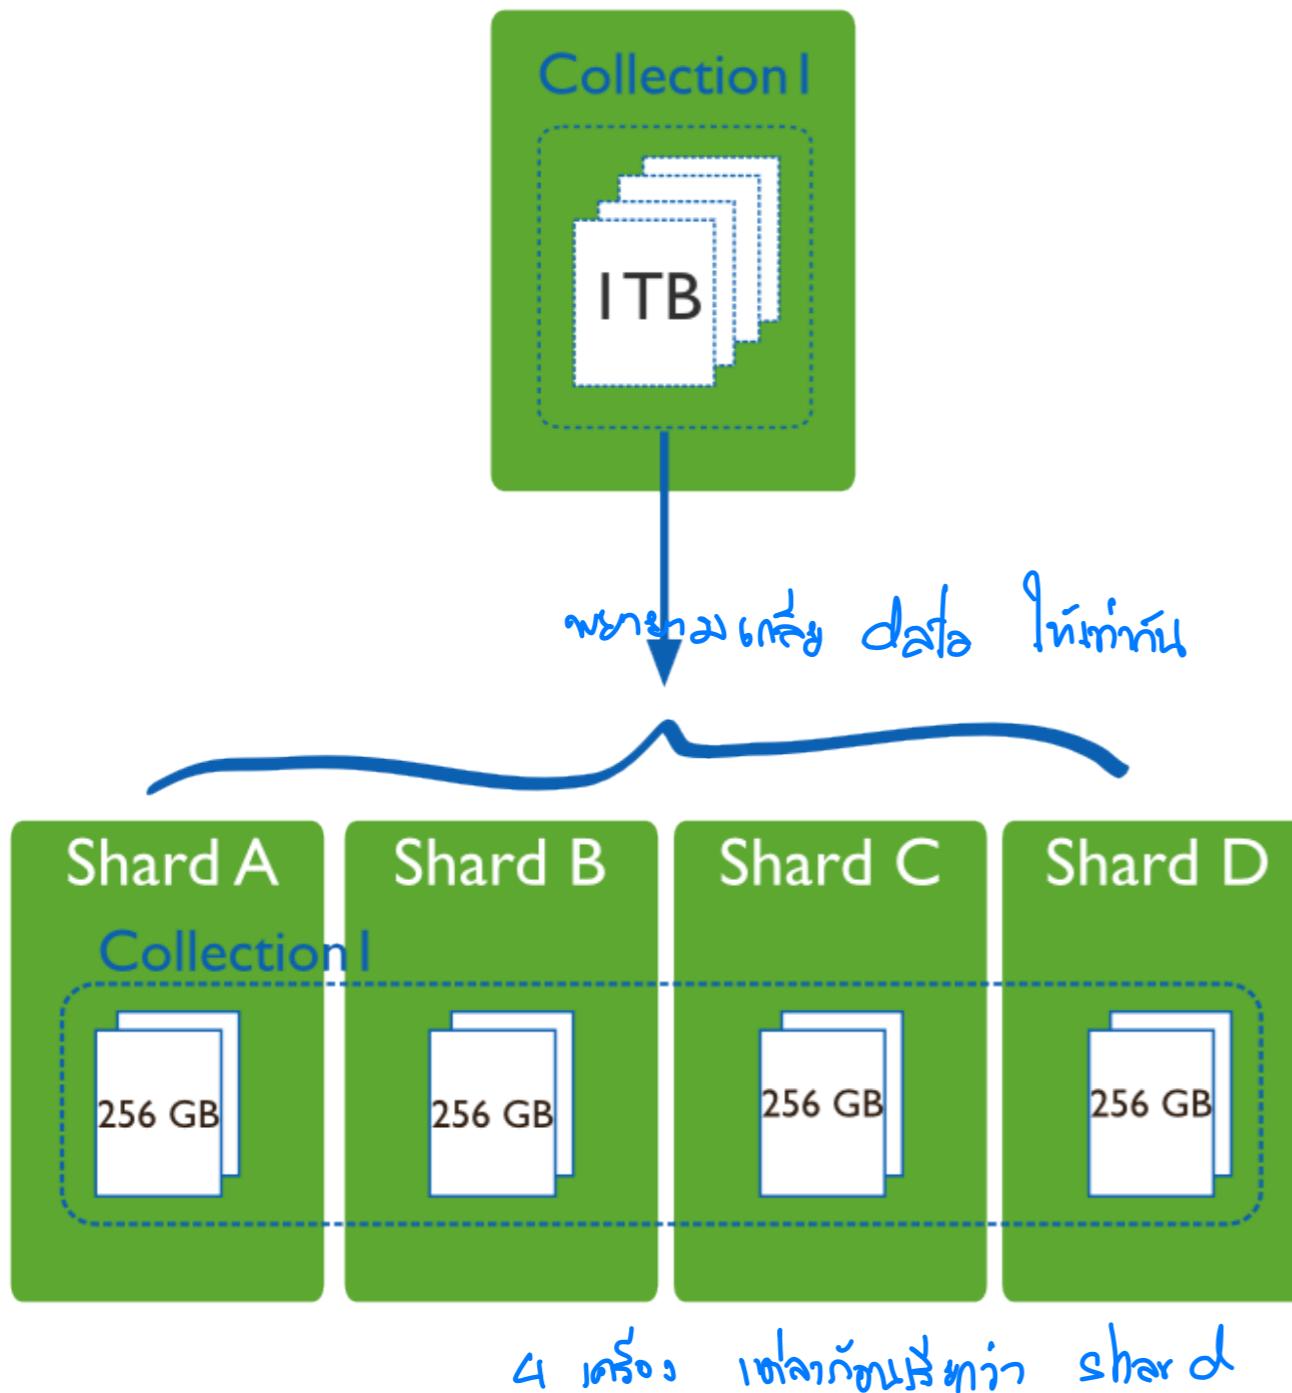
Playing with MongoDB Replica Set

- Start all containers
docker-compose up
- Initialize the replica set (inform just one node)

```
docker exec -it mongo_n1 bash -c
"echo 'rs.initiate({_id : \"rs1\",
members: [{ _id : 0, host :
\"mongo_n1\" }, { _id : 1, host :
\"mongo_n2\" }, { _id : 2, host :
\"mongo_n3\" } ] })' | mongo"
```

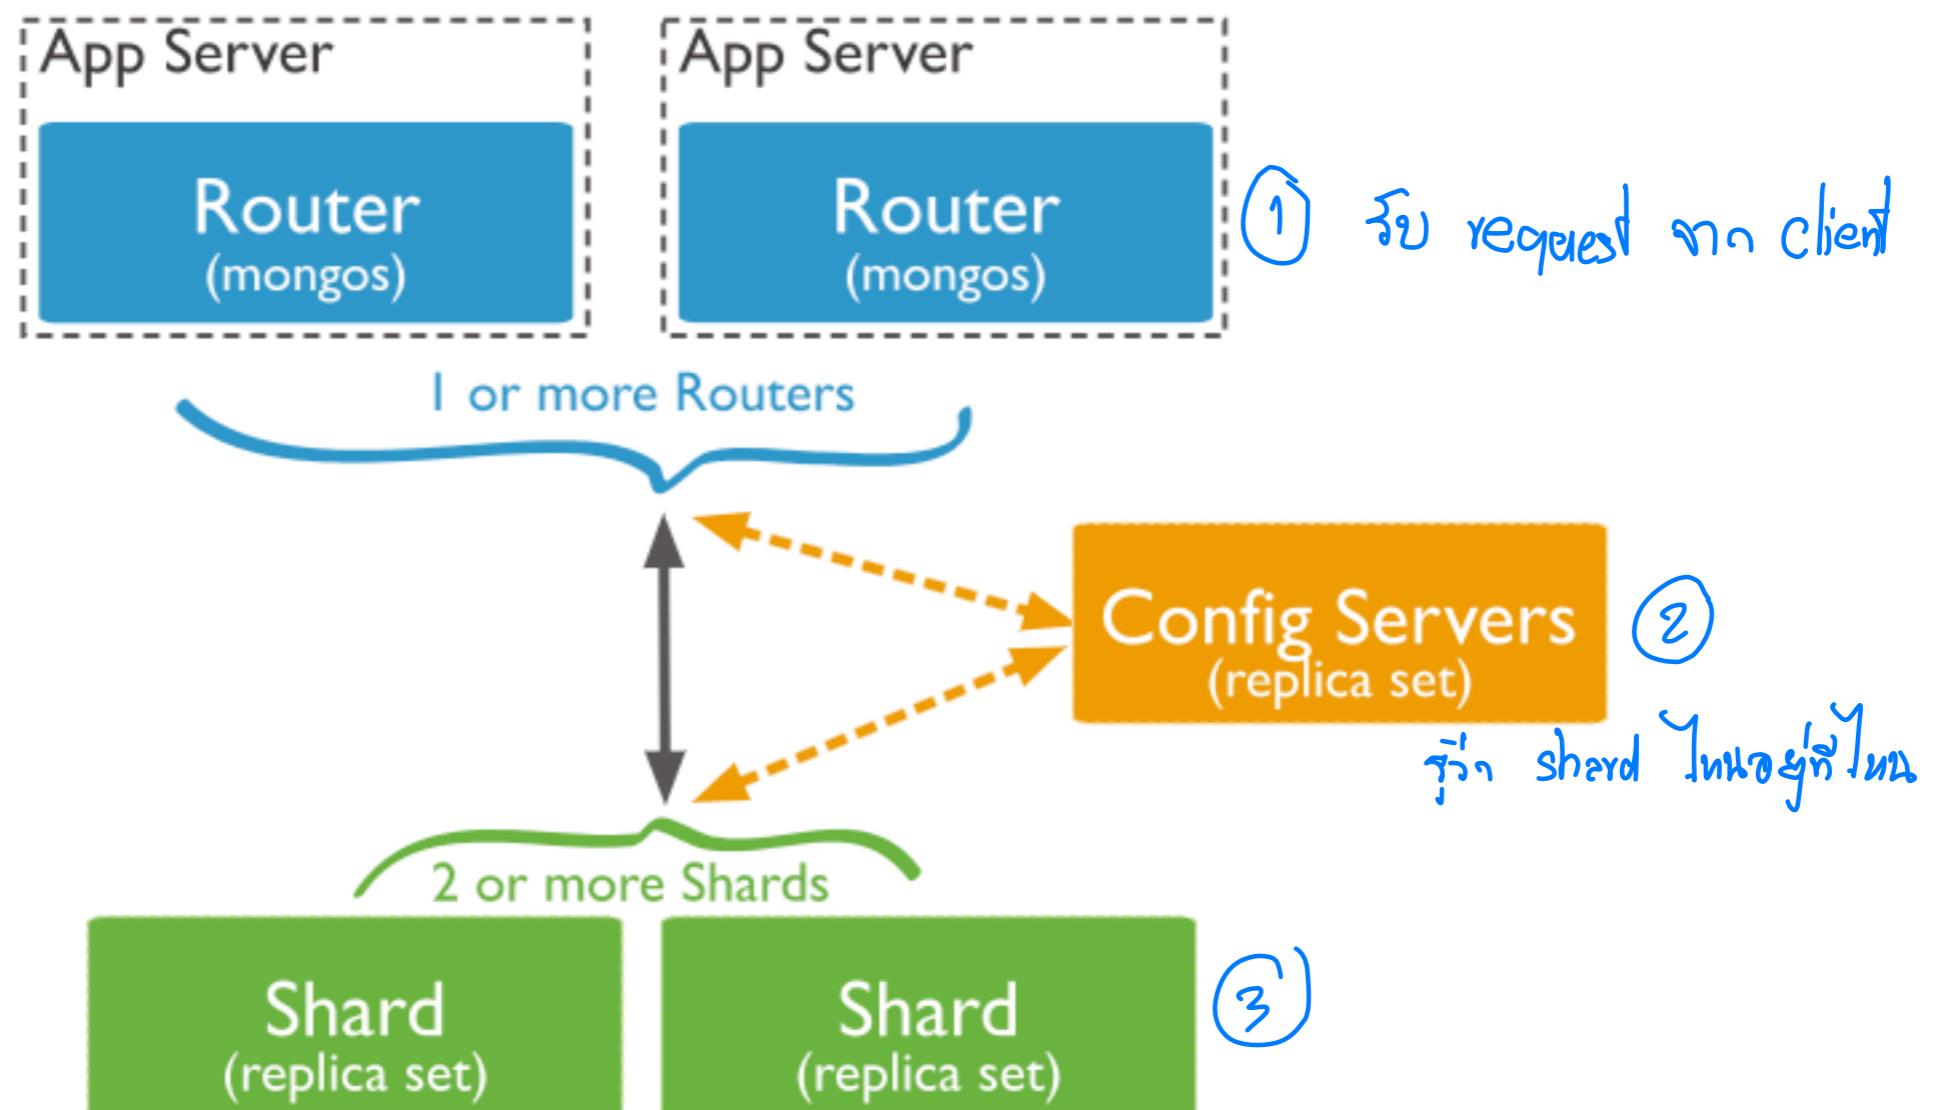


Scalable with Sharding

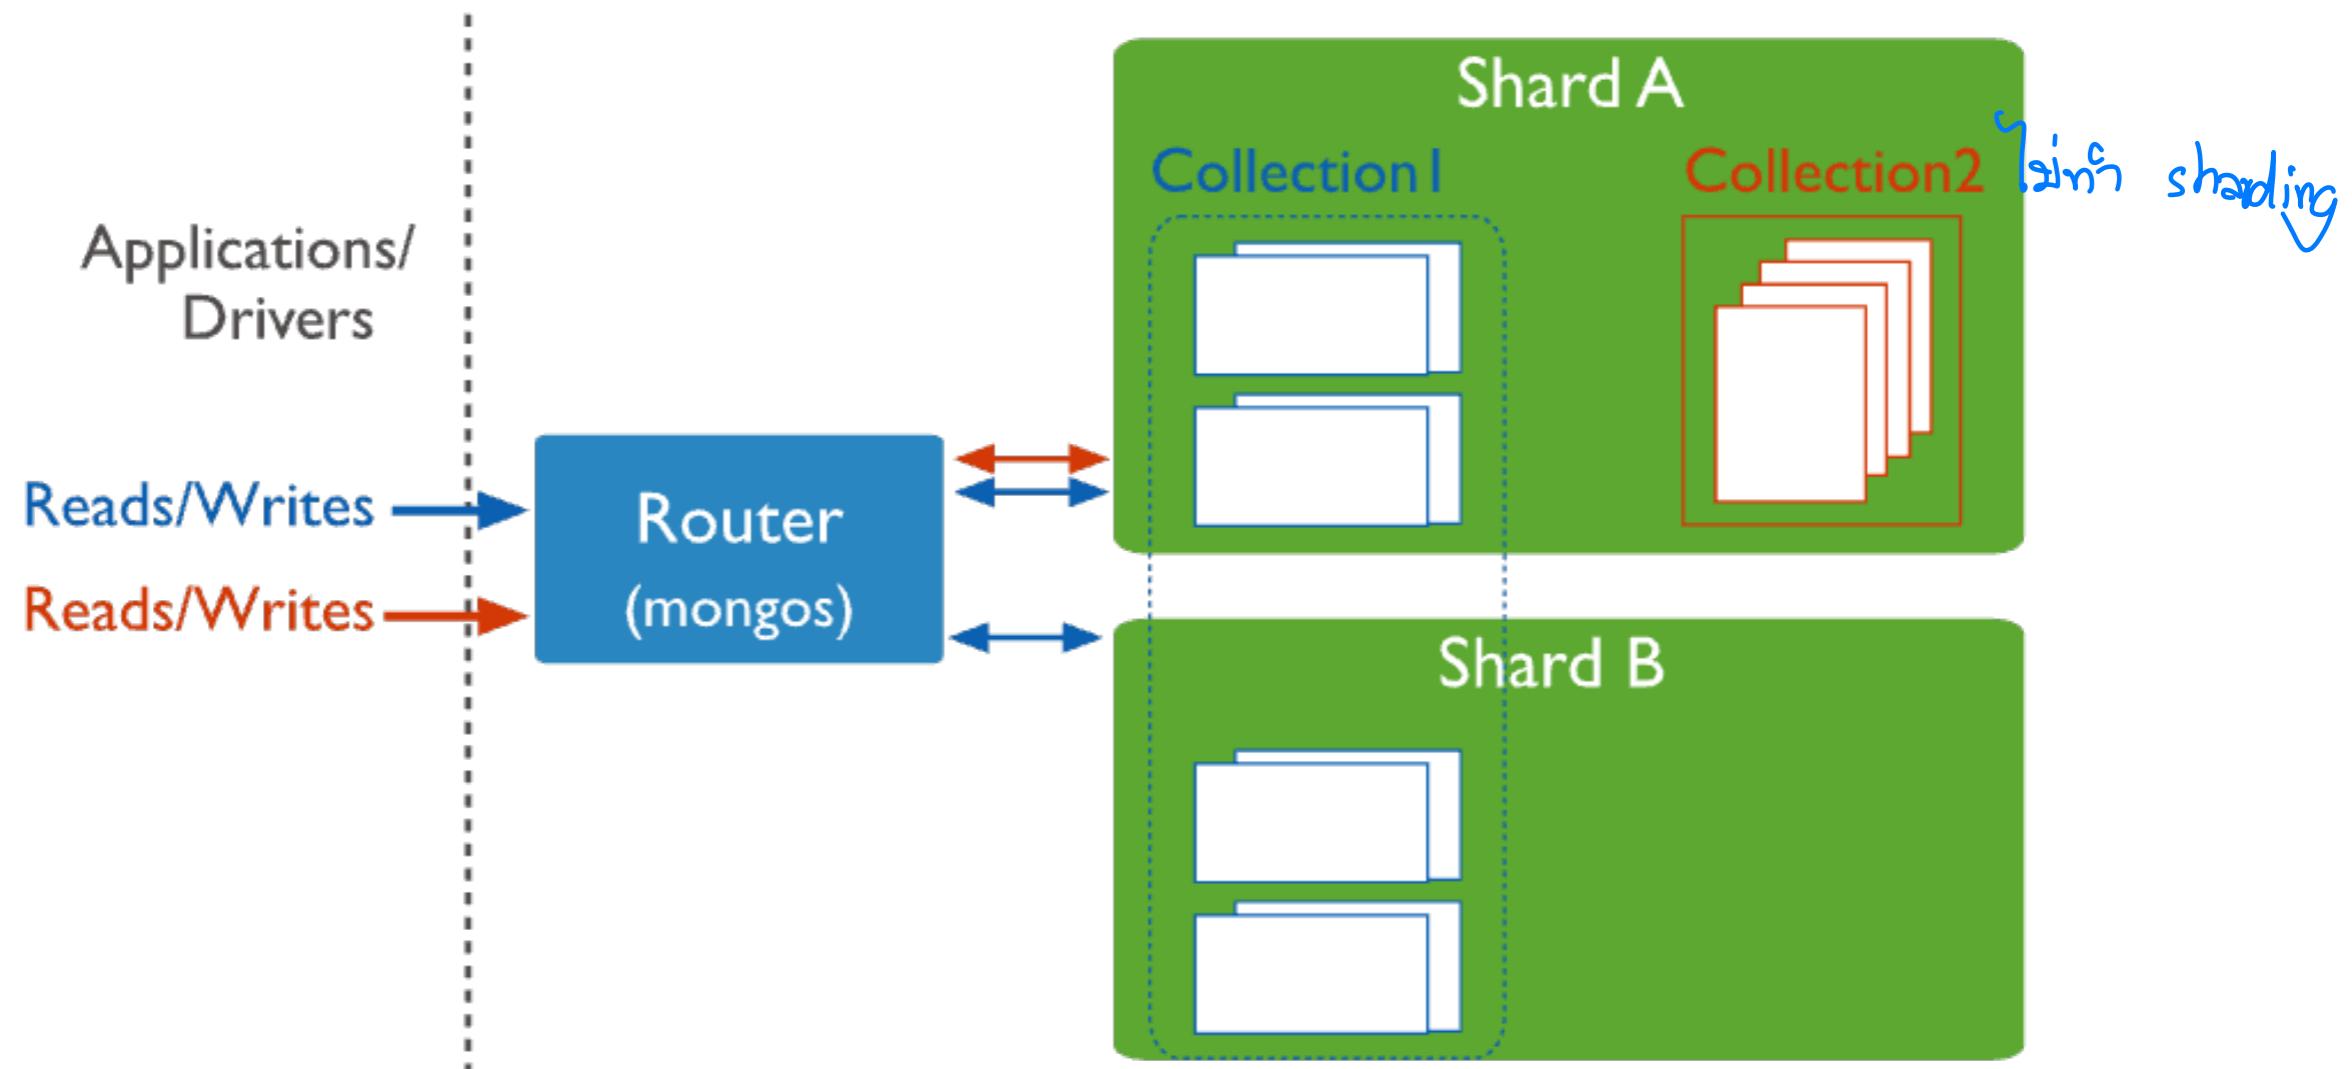


- Large-Scale MongoDB bases on Horizontal Scaling mechanism
- Storing shards (blocks of data) across multiple machines
- Data partitioned into shards with shard keys
- Each shard handles only operations related to its block

Shard Cluster Deployment



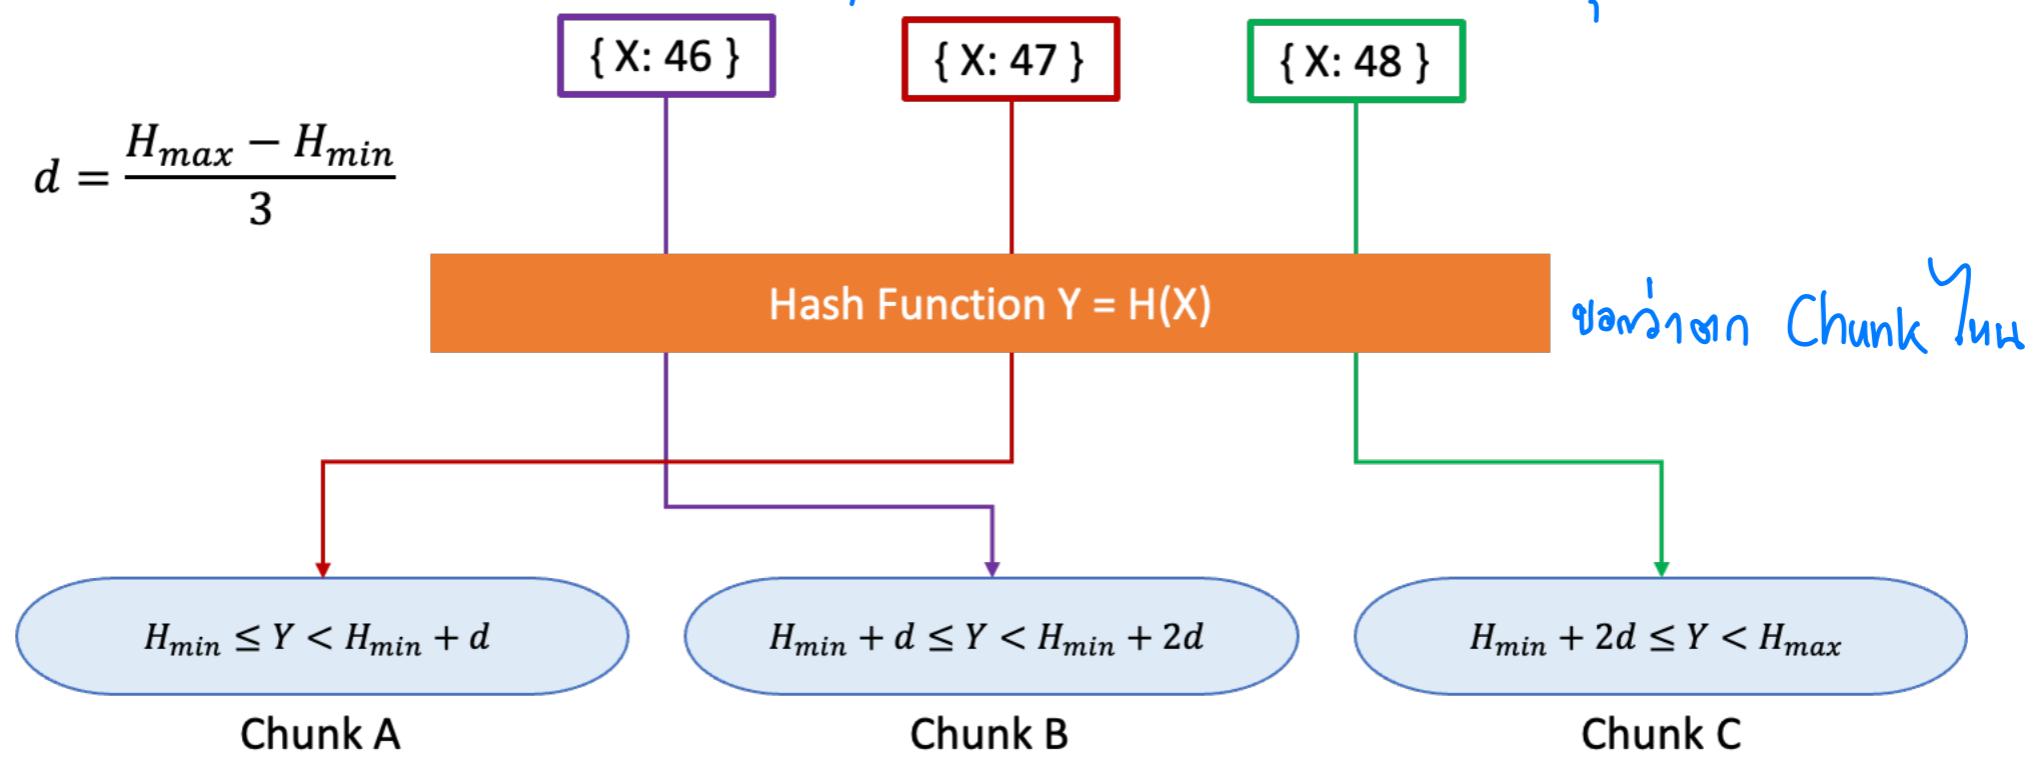
MongoDB Sharding



Shard Key Concepts: Hash Sharding

ចំណាំ: ក្នុងវិវាទ load ត្រូវឱ្យបានលើកនឹងរឿង

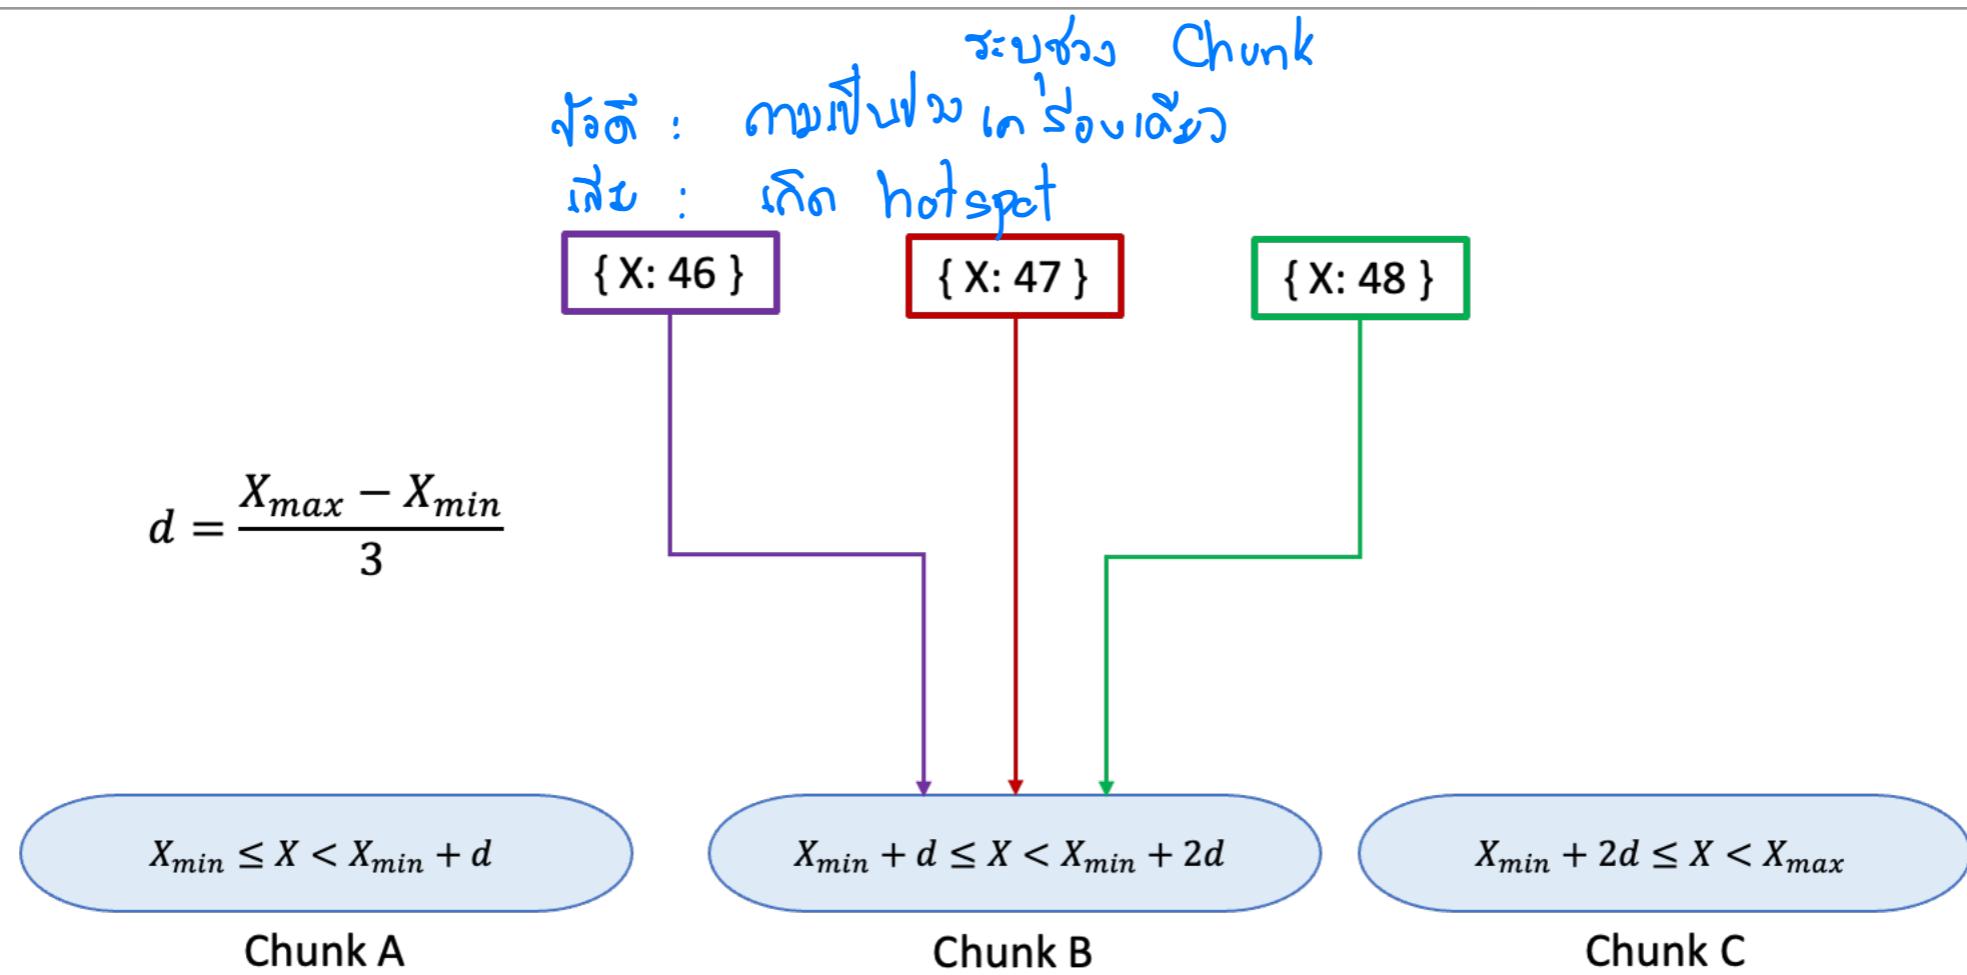
ចំណាំ: តាត query ដែលមាន (0-50) ព័ត៌មានអាជីវកម្ម



និត្យនា random access រាយការណ៍ ដៃវាទេរងទីតួកគ្នា

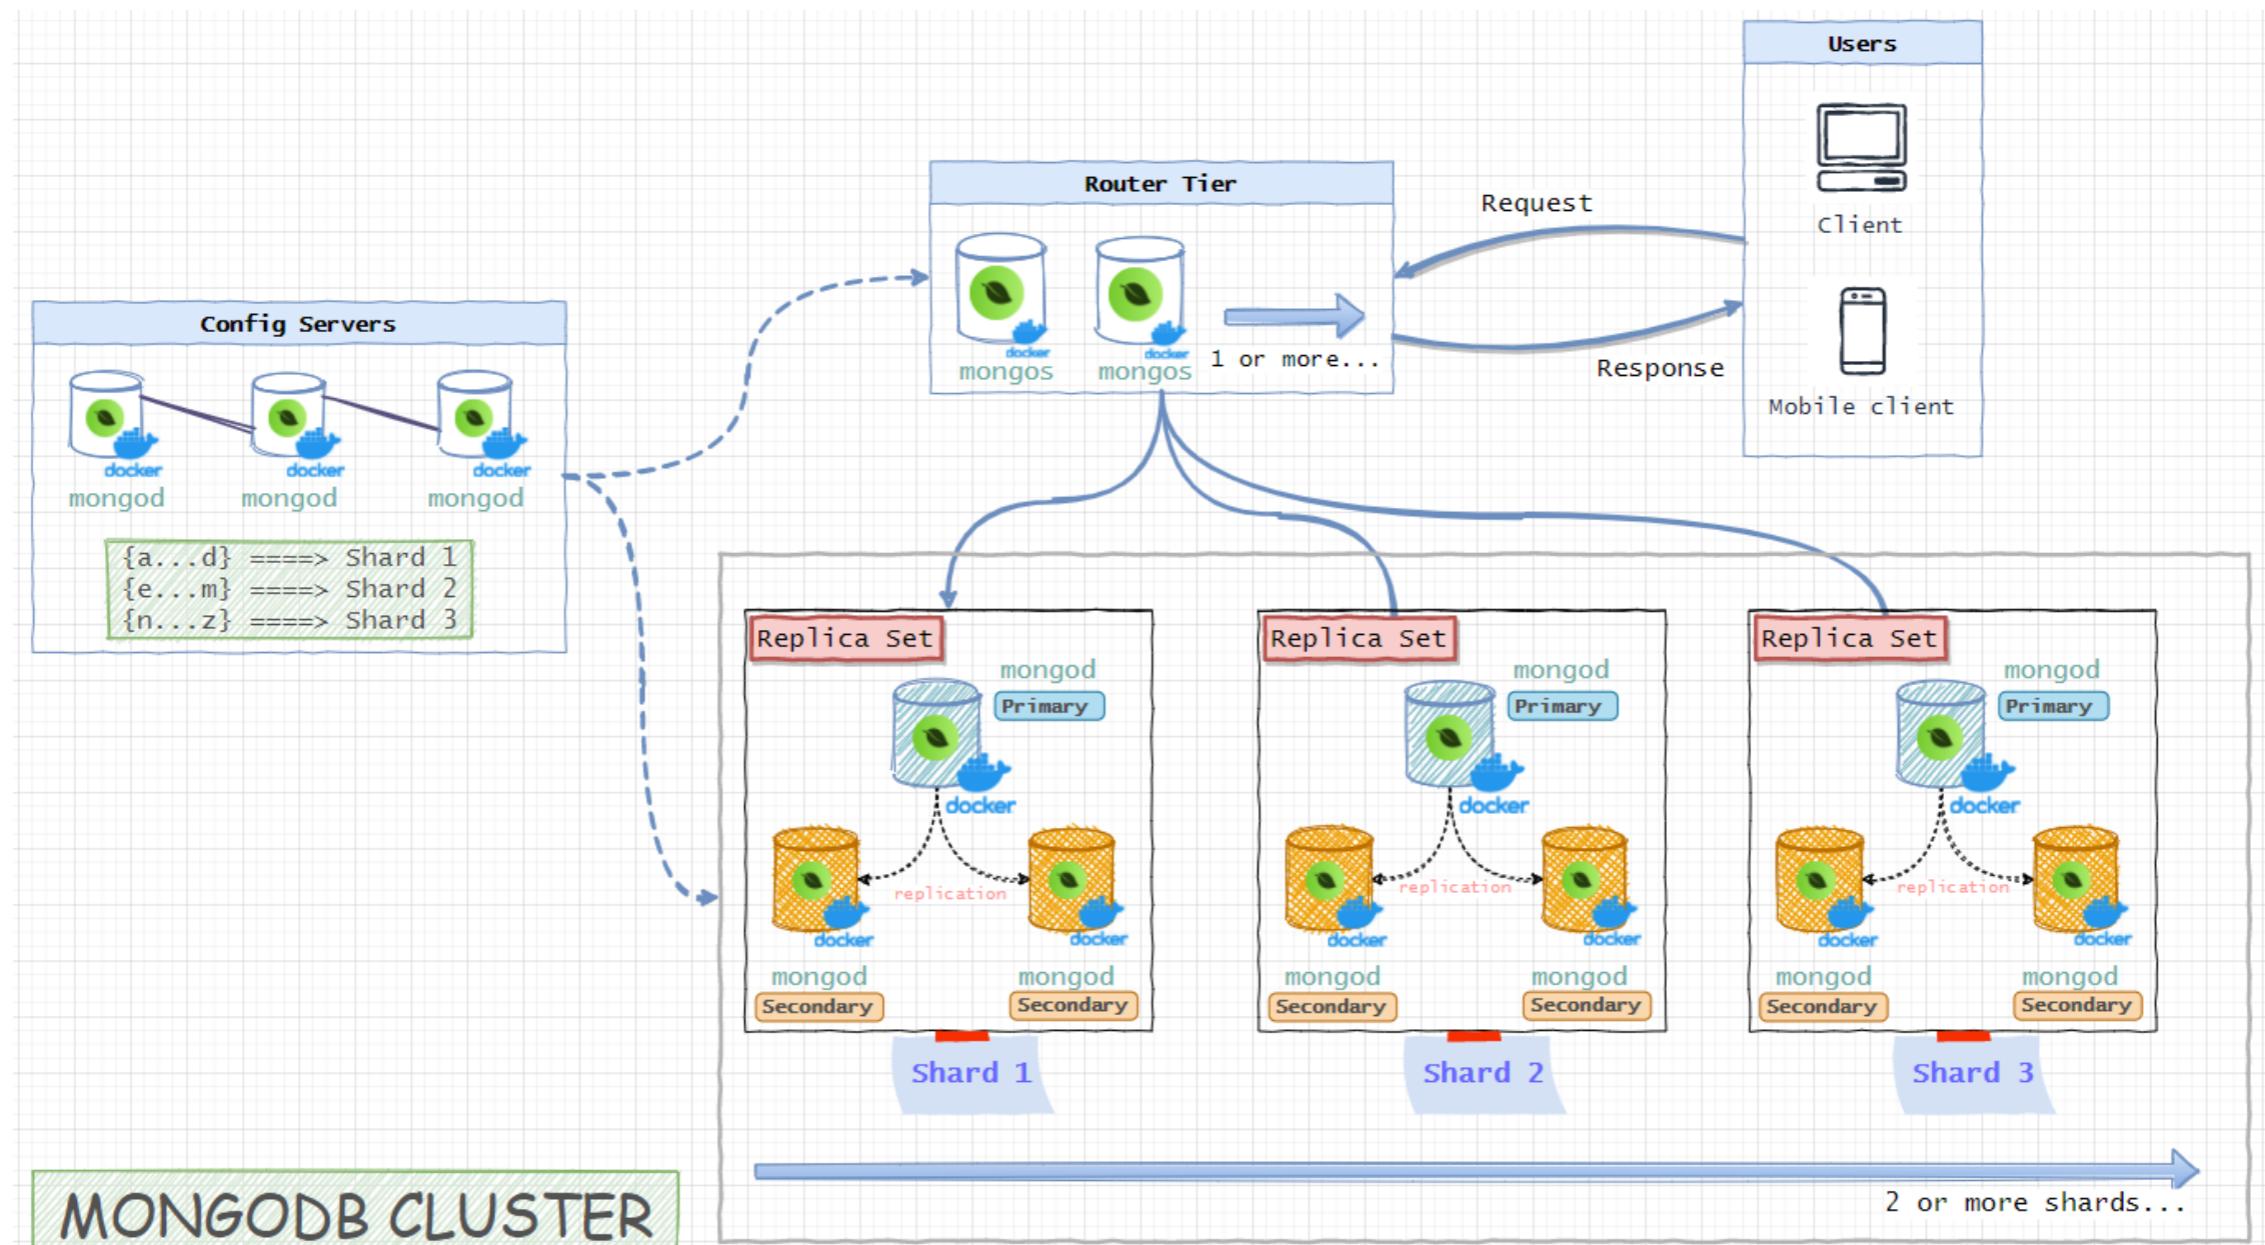
ឱ្យដឹង hotspot (R/W ទូទាត់ទាញឱ្យក្នុង)

Shard Key Concepts: Range Shading



Playing with Shards

- Refer to <https://github.com/minhhungit/mongodb-cluster-docker-compose>



MONGODB CLUSTER

<https://github.com/minhhungit/mongodb-cluster-docker-compose>



DynamoDB



Mongo relational RDBMS
Dynamo non performance cloud service

Document Store
Amazon DynamoDB

Amazon DynamoDB

- Managed NoSQL Database Services on Amazon Web Services (also available in On-Premise)
- Document-oriented and key-value models
- (Somewhat) flexible schema - Table / Items / Attributes
- Highly available with 3x replication
ถ้าต้องการให้คงรักษาไว้ต้องดูแลอย่างดี
- Trade off scalability with relax data consistency (AP NoSQL)
- Very fast, milliseconds latency at any scale
- Utilize Solid State Drive
ใช้ replicate หาข้อมูลที่ต้องใช้ล่าสุดทุกครั้ง
- Push consistency toward read operations
กำหนดให้ replication เก็บเวลาที่มีผลลัพธ์ทันท่วงทัน
- Support several language bindings:
Java, Node.js, C# .NET, Perl, PHP, Python, Ruby, Haskell and Erlang

RDBMS vs. DynamoDB

RDBMS

Database

Table

Row / Tuple

Column

Table Join

Primary Key

~~MongoDB~~ *DynamoDB*

Database

Table

Item

Attribute

Embedded Item

Partition Key + Sort Key

Dynamo កំនើងថ្មី document store នឹង

DynamoDB Data Model

Partition Key + Sort Key = must unique Table

for search in node/partition នេះ
for search in node/partition ត្រូវ sort data តុងក្នុង node
ការសរសៃការ indexing

Item	Partition Key	Sort Key	attribute1 = value	attribute2 = value	attribute4 = value	
Item	Partition Key 1	Sort Key 1	attribute1 = value	attribute2 = value	attribute4 = value	
	Partition Key 1	Sort Key 2	attribute1 = value			
	Partition Key 2	Sort Key 5	attribute2 = value	attribute3 = value	attribute4 = value	attribute8 = value
	Partition Key 3	Sort Key 5	attribute1 = value	attribute5 = value		

តាម query សម្រាប់ក្នុង category

សៅរ៍ទាំងអស់ by name និងចុច្ចបញ្ជាក់របស់ខ្លួន

category	name	cuisine	rating	votes	(list of) ingredients
category	name	cuisine	rating	votes	(list of) ingredients

រួមទាំង
design នាម
query

name	category	cuisine	rating	votes	(list of) ingredients
name	category	cuisine	rating	votes	(list of) ingredients

DynamoDB API

CreateTable

GetItem

UpdateTable

Query

DeleteTable

Scan

DescribeTable

BatchGetItem

ListTables

PutItem

UpdateItem

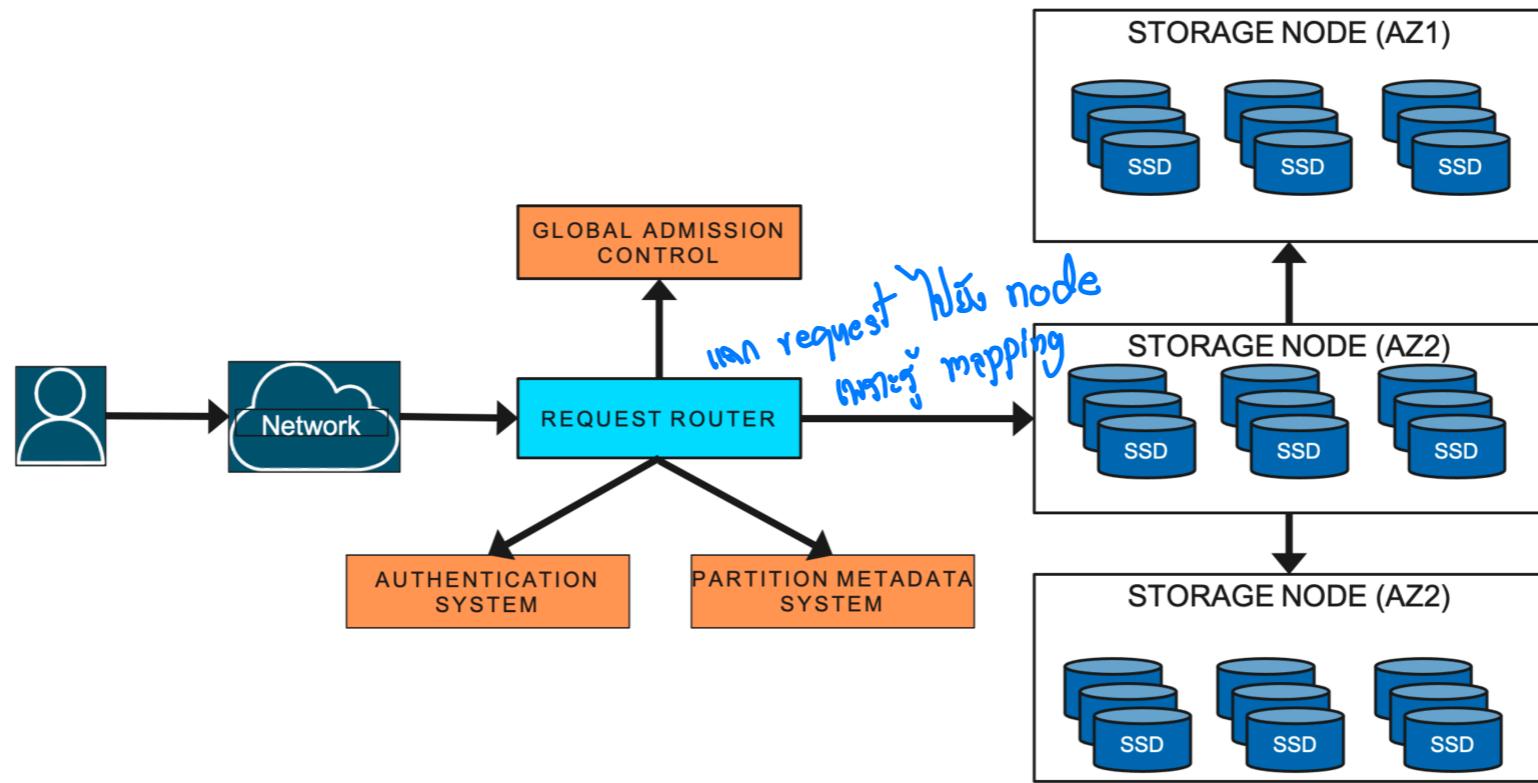
DeleteItem

BatchWriteItem

Data Types for Attribute

- String (S)
- Number (N)
- Binary (B)
- String Set (SS)
- Number Set (NS)
- Binary Set (BS)
- Boolean (BOOL)

DynamoDB Architecture



- Items are distributed among partitions resided in storage nodes using hashed Partition Key (aka. Hash Key)
- At each storage node, items are stored in B-Tree structure stored by Sort Key

DynamoDB Partition

- Partitions host disjoint and contiguous parts of the table's key-range
- Each partition can store up to 10GB with 3000 RCUs and 1000 WCUs; partition is splitter if exceeding these numbers
- Each partition has multiple replicas distributed across different Availability Zones for high availability and durability
- The replicas for a partition form a replication group and one replica in the group is elected as the leader (or primary)
 - With write request, leader replica generates a **write-ahead log** record and sends it to its peer (replicas) and returns a successful response to the client when **one of the replicas** responds with a successful write — eventually consistent write ($W=2$)
 - With read request, client can choose either **strongly consistent read** ($R=N-1$) and **eventually consistent read** ($R=1$)

Playing with Amazon DynamoDB Locally

- Amazon provides DynamoDB docker container to try locally

```
docker run -d -p 8000:8000 --name dynamo --rm amazon/dynamodb-local
```
- You can access local DynamoDB services using aws CLI or Python with boto3 package
- Note: you will need aws access credential including AWS Access Key ID and AWS Secret Access Key
 - create them from AWS console and config to your machine with **aws configure** command

Dynamo និង semi - schemas

Why Document Oriented NoSQL?

អាជីវកម្ម partition + shard key

- “Natural-fit” to lots of information
 - Parent-child objects (e.g. story-chapters-comments)
- Schema changed over time *schema evolution*
- Expecting high write load
- Horizontal scaling is needed in the future (e.g. data will grow very large) *scaling data ក្នុងបច្ចេកទេស (ធ្វើជាន់ឡើង)*
- Data backup/restore with point-in-time approach

Assignment Readings

- M. Elhemali, et al. "Amazon DynamoDB: A Scalable, Predictably Performant, and Fully Managed NoSQL Database Service", 2022 USENIX Annual Technical Conference (USENIX ATC 22). 2022.

References

- Nic Raboy, “Getting Started With MongoDB As A Docker Container Deployment”, <https://www.thepolyglotdeveloper.com/2019/01/getting-started-mongodb-docker-container-deployment/>
- อภิศิลป์ ตรุษกานนท์, “ทำไม Pantip ปิดปรับปรุงบ่อย บทเรียน ความเจ็บปวด และคราบน้ำตาจาก MongoDB”, <http://macroart.net/2013/10/mongodb-lessons-learned-on-pantip/> เกี่ยวกับเรื่อง sharding ณ ณ
- T. Singthong, “แนวทางการสร้าง Index และ การเลือก Sharding strategy ใน MongoDB”, <https://medium.com/iamgoangle/แนวทางการสร้าง-index-และการเลือก-sharding-strategy-ใน-mongodb-d370408df0f6>
- T. Gamage, “Running AWS DynamoDB Local with Docker-Compose”, <https://medium.com/platform-engineer/running-aws-dynamodb-local-with-docker-compose-6f75850aba1e>