



2110625 - Data Science Architecture

Key-Value Stores

Asst.Prof. Natawut Nupairoj, Ph.D.
Department of Computer Engineering
Chulalongkorn University
natawut.n@chula.ac.th

Cassandra កំណើន memory នៅក្នុងរឿង
នៃ disk និង recovery ។

Dynamo នៅក្នុងទីនាំ memory

Key-Value Stores

ក្នុងគម្រោងនេះ key value នៅក្នុងទីនាំនៃ memory នូវឯកសារ storage នៅក្នុង primary នៃ memory

- A data store designed for storing, retrieving, and managing associative arrays (aka. dictionary or hash table)
- Main concept is to store data as a collection of key-value pairs in which a key serves as a unique identifier
- Simple and fast, often use in-memory architecture, ability to scale-out
 - កំណើន local Storage នៃ key:value store ។
- Use cases: session store, shopping cart, SQL/API caching, etc.

↳ ក្នុងការ call API និង overhead ព័ត៌មានអាជីវកម្មនៃ query ត្រូវបាន

db သဲ cache အနေဖြင့် အခြေခံပွဲများ db မျှတိုက် app server ရှိခိုးပါမ်းများ

Redis

Key-Value Stores

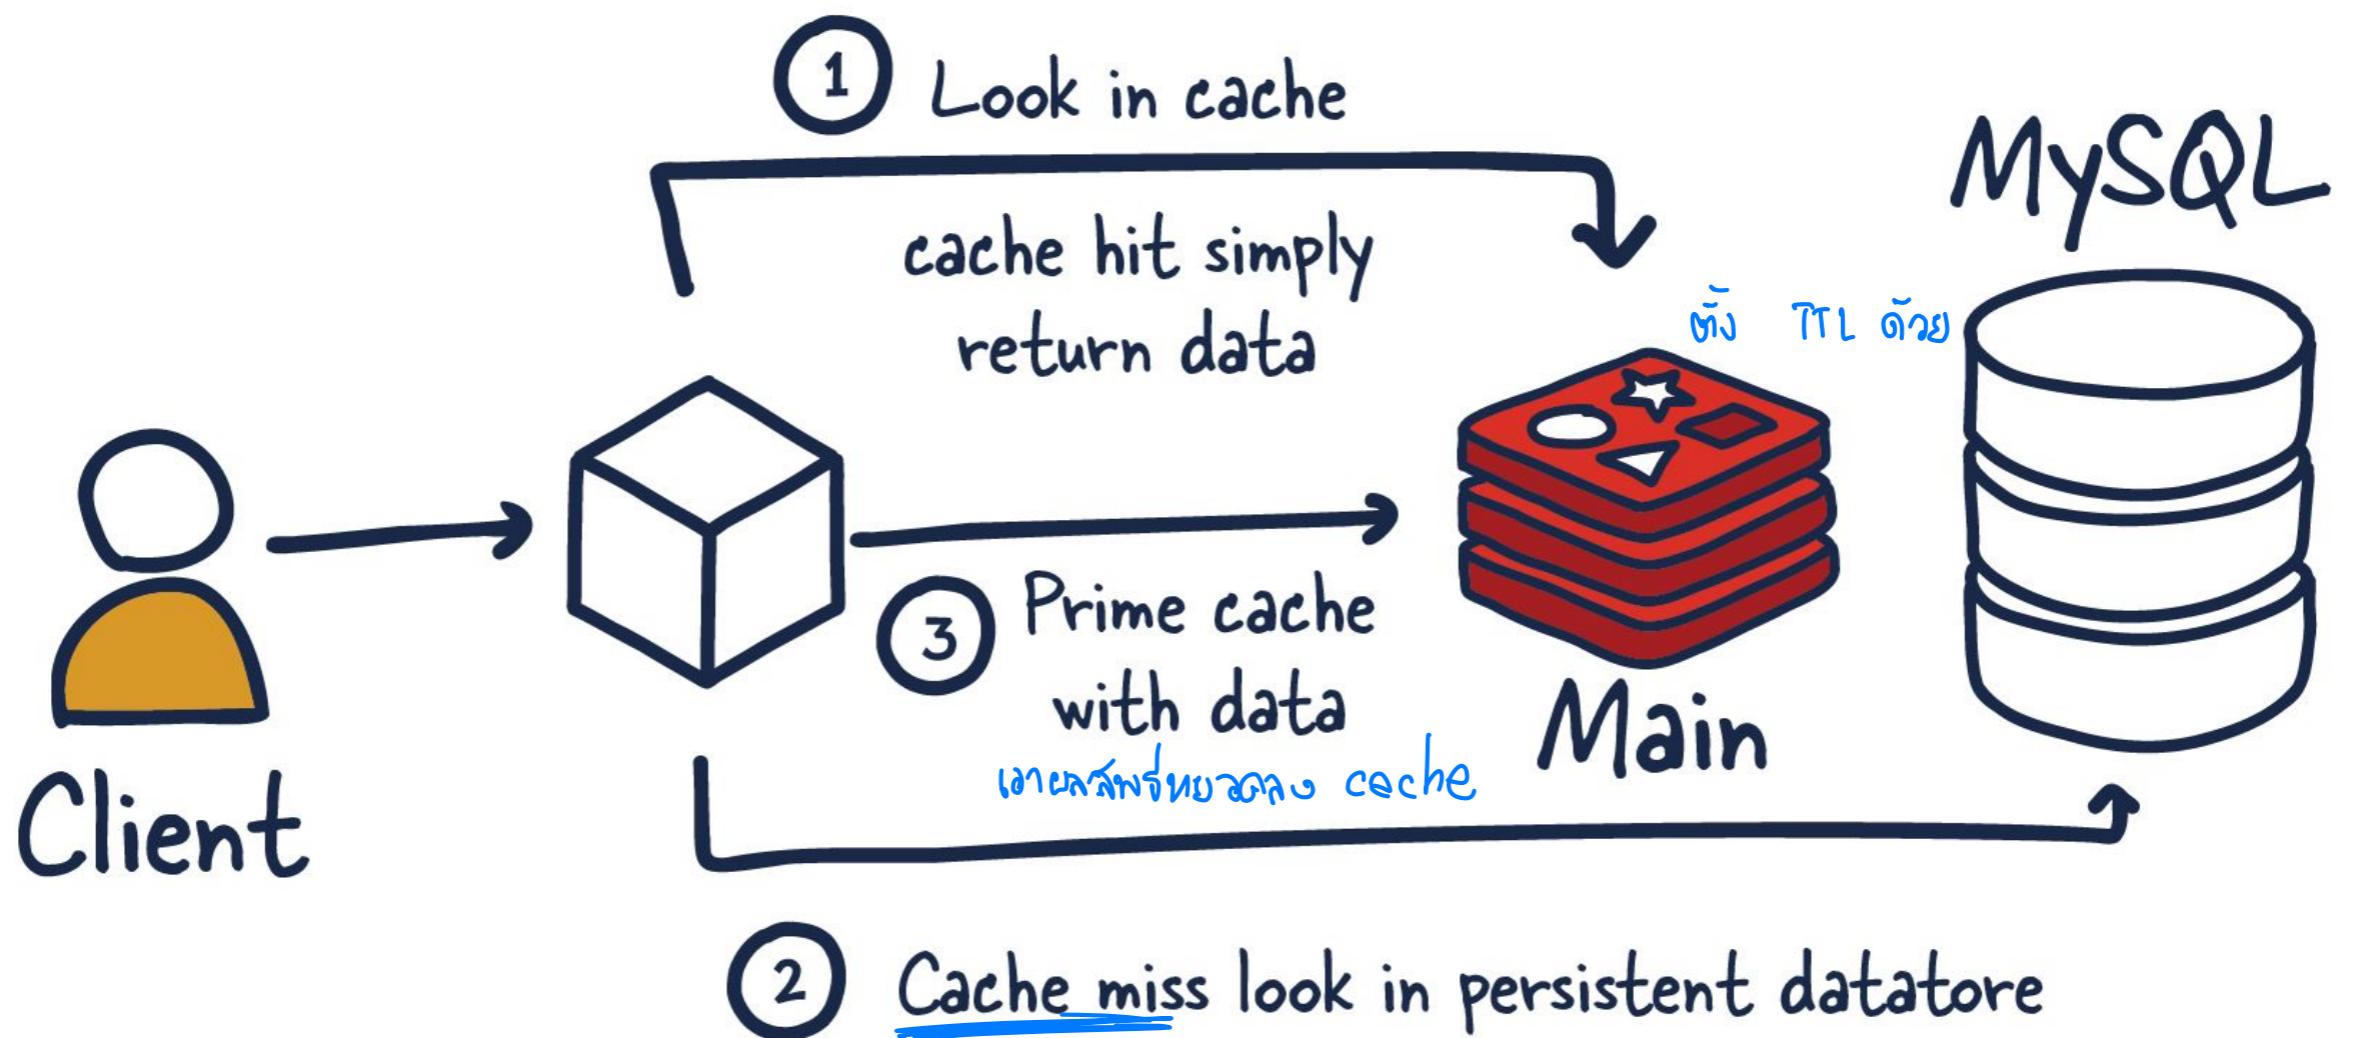


Redis (Remote Dictionary Server)

↑ ពីនិង failover ស្ថិត នៅក្នុង

- In-memory data structure store with clustering, transactional, time-to-live limiting, and auto-failover capabilities
 - key និង expire នឹងវិនិច្ឆ័យ
- Support wide-range of data structure with lots of related operations for each structure
 - រាយការណ៍បច្ចុប្បន្ន
- Being used for database cache, message broker, streaming engine, feature store engine, etc.
- Provide CLI and support many programming languages
- Many useful modules are available to extend the functionality of Redis core e.g. RedisJSON, RedisSearch, RedisTimeSeries, RedisOM, etc.
 - feature store និង dataset ដើម្បី train model ។

How is redis traditionally used



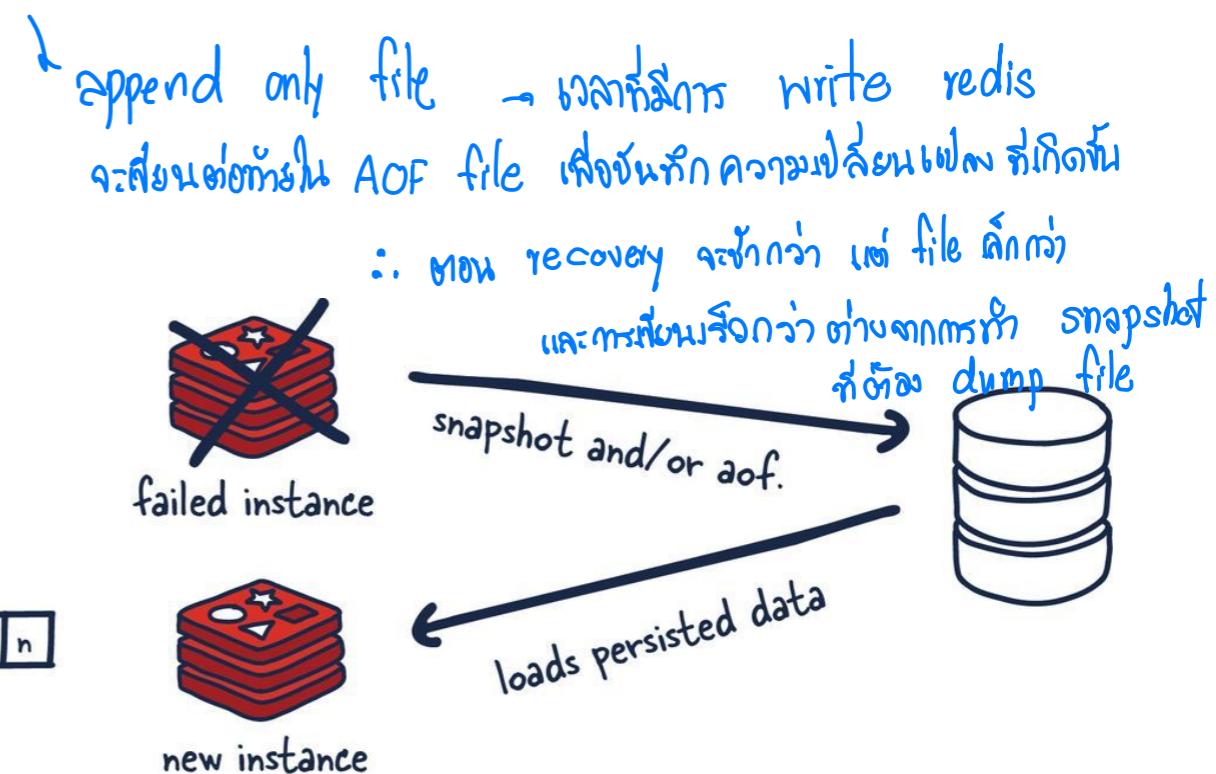
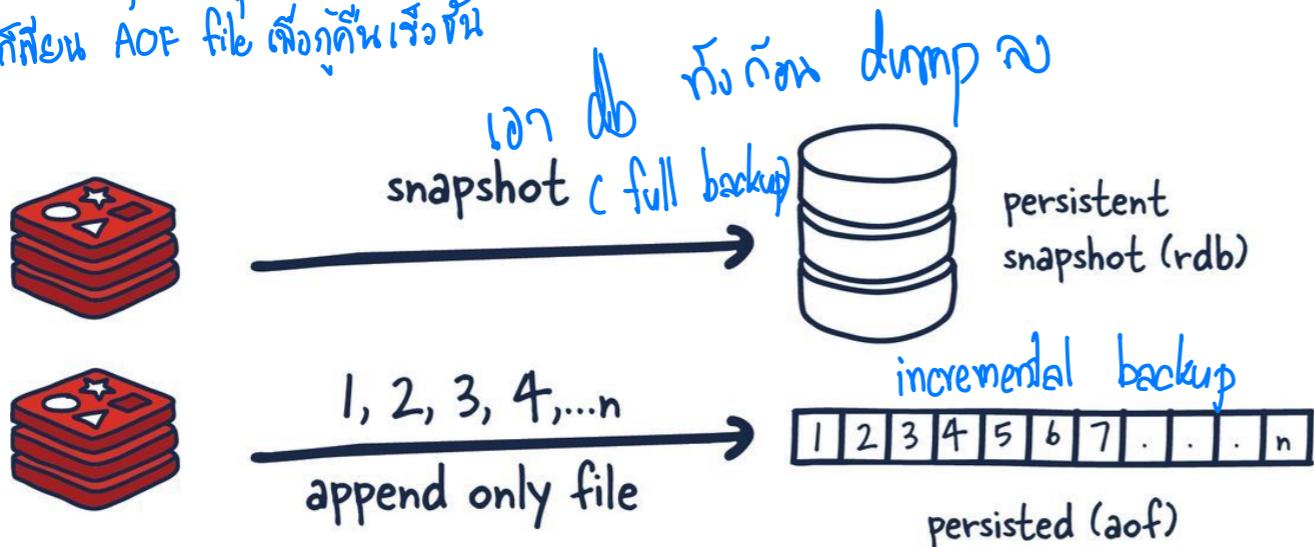
ด้วย memory เท่านั้น ก็สามารถรักษาได้

Redis Persistence

- Redis supports many level of persistence: no persistence, RDB (point-in-time snapshot), AOF (log every write),

RDB+AOF

ทำ snapshot ทุก 1 ชม. เพื่อรักษาไว้ชั่วโมง
เก็บยัง AOF file เดิมๆ คืนเร็วที่สุด



snapshot 1 ชม. ระหว่างปั๊มน้ำ AOF

redis ทำการ compress ต่อไป

Running Redis

- The simplest way to run a redis instance is to use docker

```
docker pull redis
docker run -d --rm --name redis -p 6379:6379 redis
```
- This will start redis in your docker at port 6379 and map the port to your localhost
- You can also use docker-compose.yml and other example files in redis folder in datasci_architecture repo

Working with Redis

- Redis CLI
 - Standard client program to connect to any redis server
 - Come with any redis installation (see: <https://redis.io/docs/getting-started/>)

redis-cli

```
redis-cli -h 34.143.227.66
```

- You can type in redis command in the CLI input

Working with Redis

- Redis-Py
 - Standard python package for redis client

```
pip install redis
```

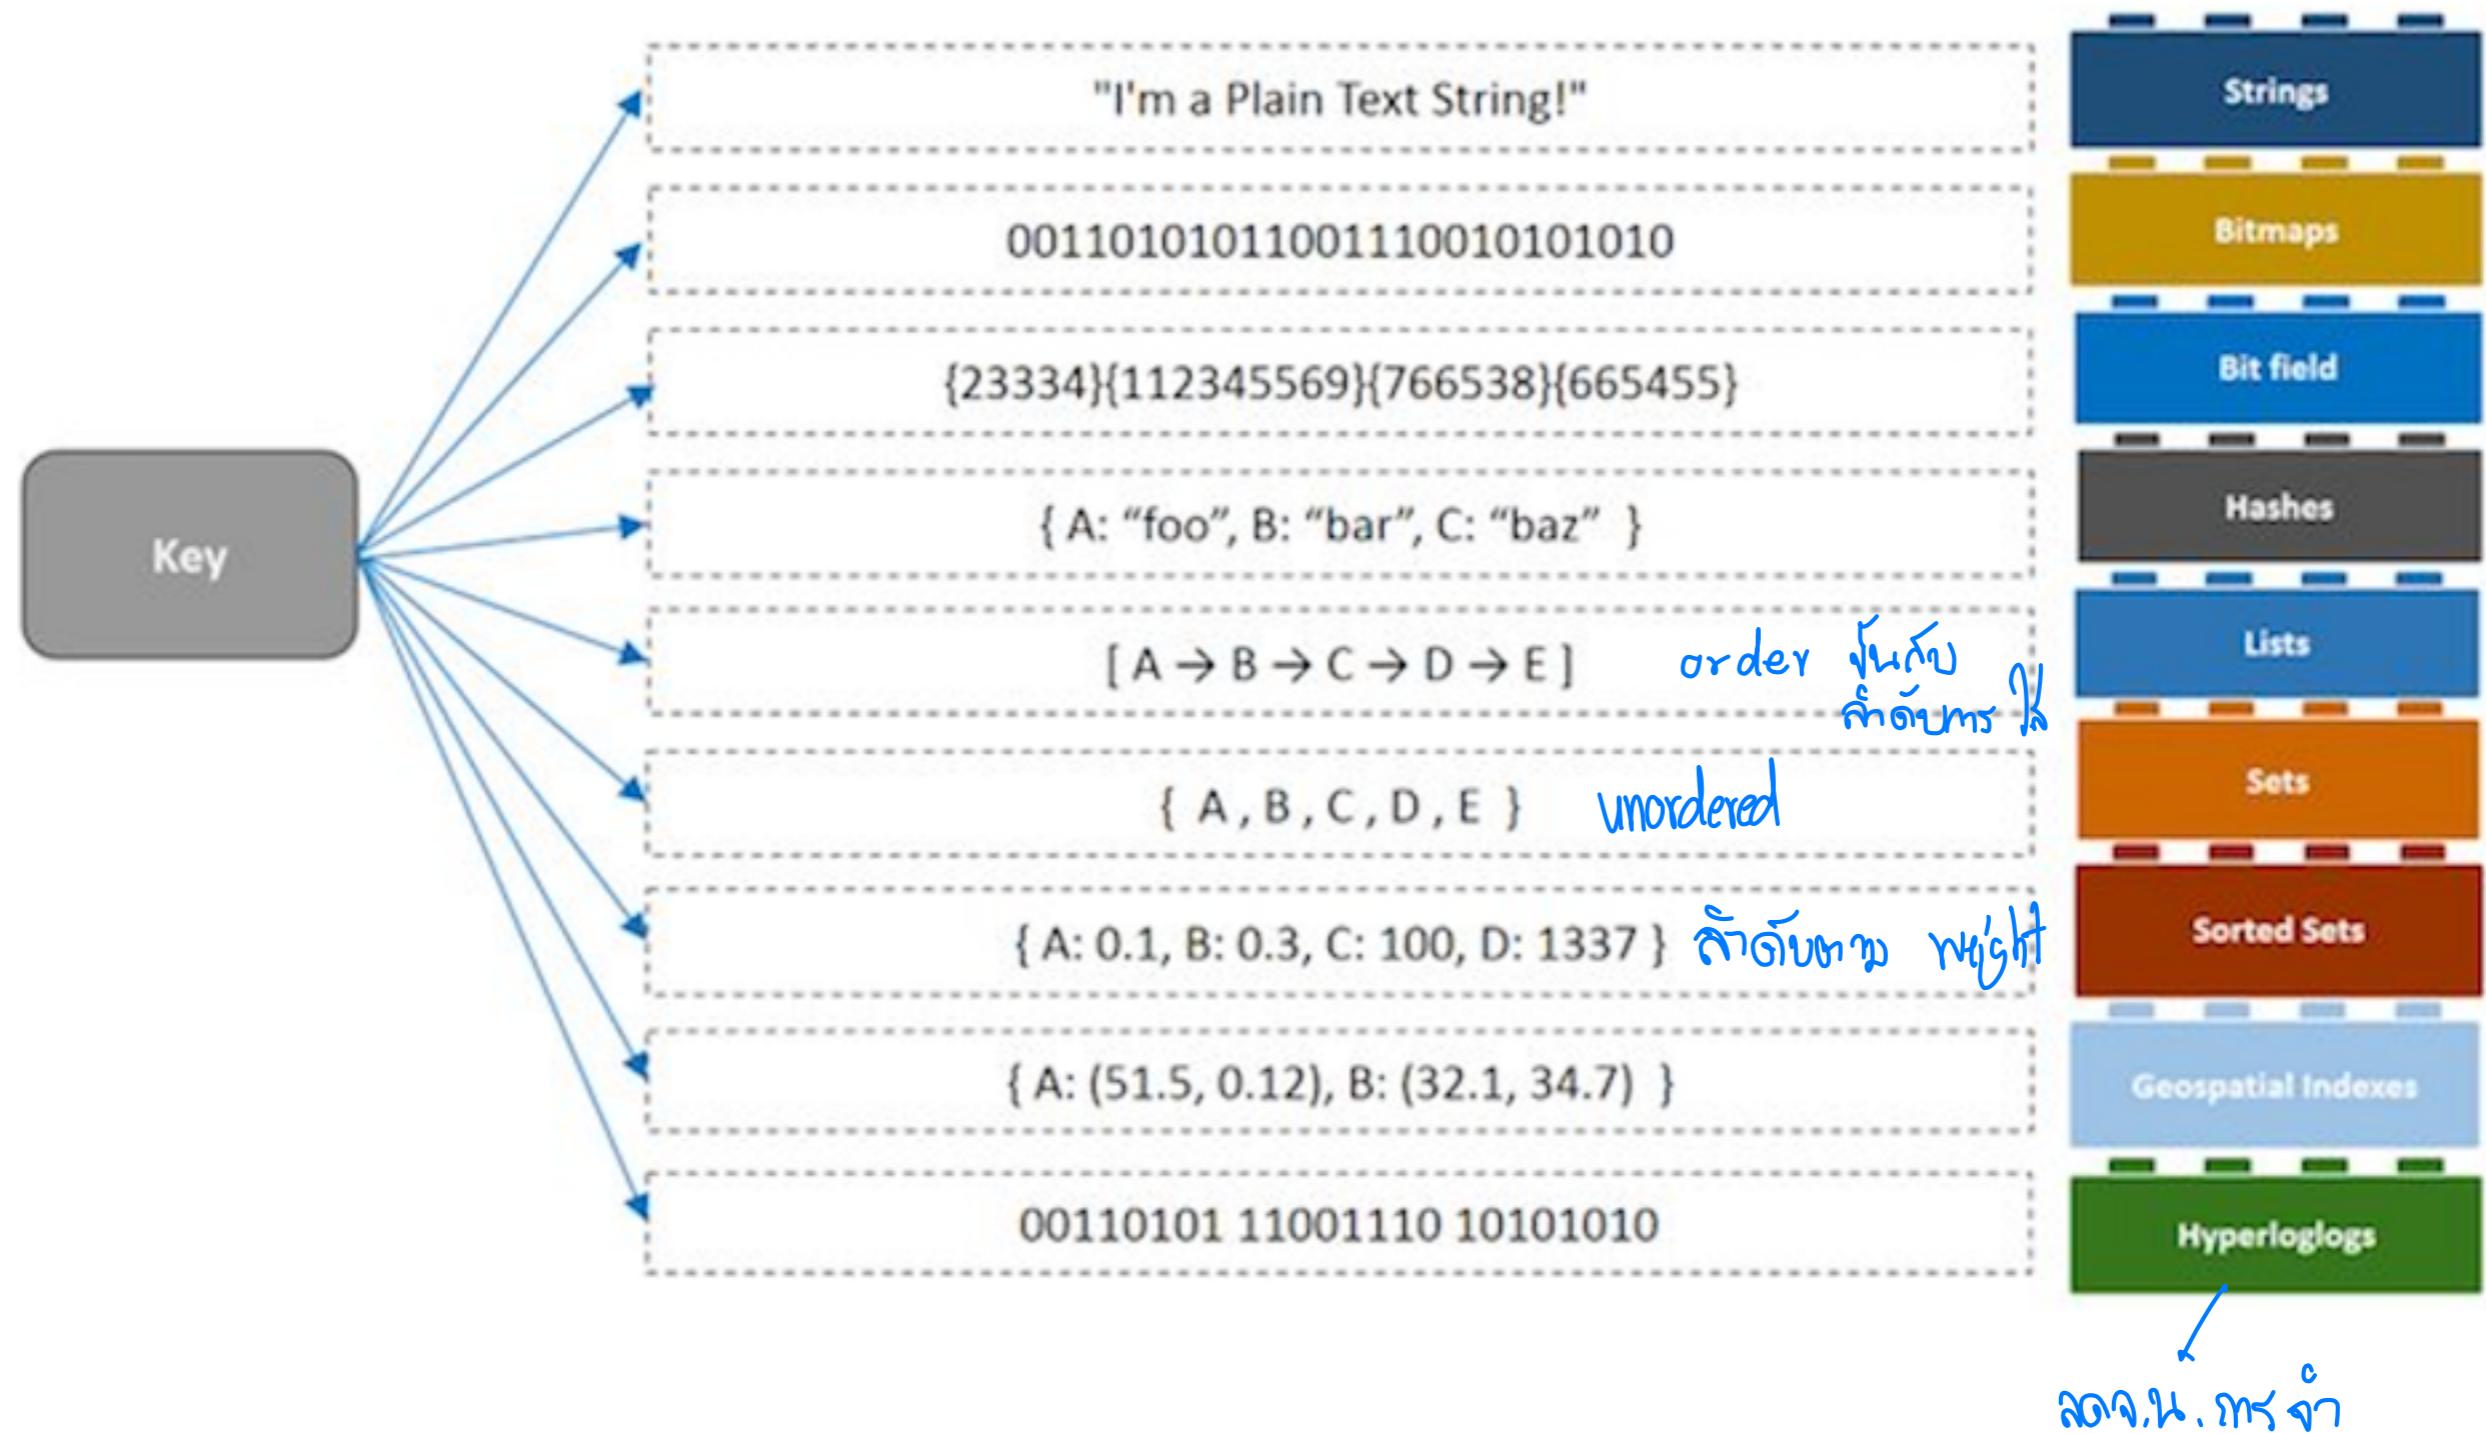
- Example

```
import redis

r = redis.Redis(host='hostname', port=port)
```

Redis Data Types

ຈຳກັບຄາມໂຕບສ້າງ value



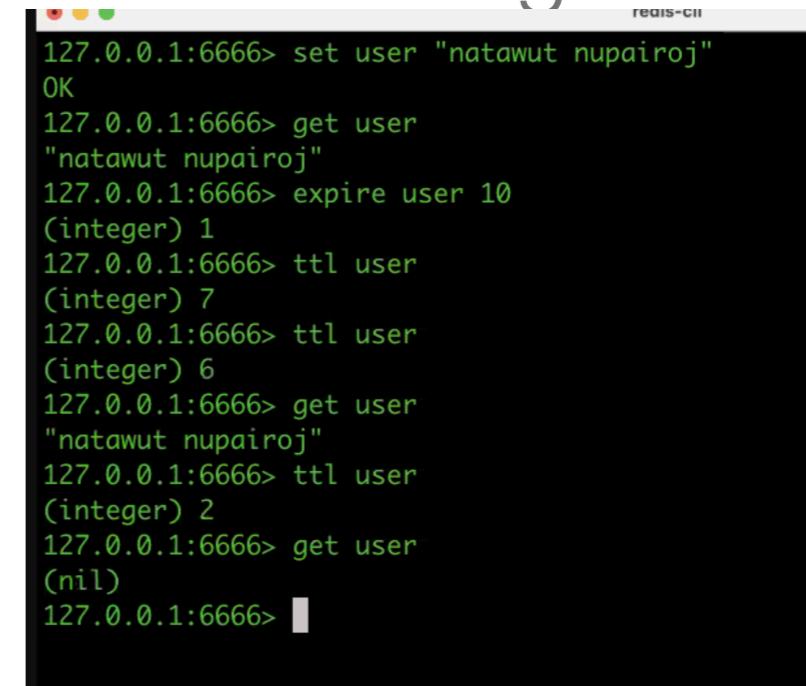
String

- Similar to Python or Java Strings, maximum length of 512MB

SET "user" "Natawut Nupairoj"
key value
GET "user"
DEL "user"

- Use cases

- Server-side object cache e.g. HTML fragments, shopping cart, user profile
- Queues
- Activity tracking



A screenshot of a terminal window titled "redis-cli". The window shows a series of Redis commands and their responses:

```
127.0.0.1:6666> set user "natawut nupairoj"
OK
127.0.0.1:6666> get user
"natawut nupairoj"
127.0.0.1:6666> expire user 10
(integer) 1
127.0.0.1:6666> ttl user
(integer) 7
127.0.0.1:6666> ttl user
(integer) 6
127.0.0.1:6666> get user
"natawut nupairoj"
127.0.0.1:6666> ttl user
(integer) 2
127.0.0.1:6666> get user
(nil)
127.0.0.1:6666>
```

Other String Commands

APPEND

INCR

SET

DECR

INCRBY

SETEX

DECRBY

INCRBYFLOAT

SETNX

GET

LCS

SETRANGE

GETDEL

MGET

STRLEN

GETEX

MSET

SUBSTR

GETRANGE

MSETNX

GETSET

PSETEX

Useful Commands

- Any item in Redis can be made to expire after or at a certain time

```
EXPIRE user 60. # in seconds
```

```
TTL user
```

- You can scan all index with scan command

```
SCAN 0
```

```
127.0.0.1:6666> scan 0 match *
1) "17"
2) 1) "run:id:8ceb7b0af5eb4836ae96e033d16edeae"
   2) "run:project:46"
   3) "project:110"
   4) "run:id:0f779b96964e4fada79ab2bf14d07f3d"
   5) "run:id:c3e47ddd9b794f56a379f779088bb910"
   6) "run:id:a49ad162270a4fe497833c4281b9a6bf"
   7) "run:id:1e25b2dece464359b8b05c1060aabb98"
   8) "run:id:892edad6c9ba47bcbbf2c6119fc7995a"
   9) "project:109"
  10) "run:id:cee505375d7f46bdad96ade14faf46f8"
  11) "run:project:109"
  12) "run:id:3389c25fe6174ec6bb4bf7683a0a99f0"
127.0.0.1:6666>
```

scan 0 = scan key őj, őm, őg

- You can delete item or test its existence

```
DEL mykey
```

```
EXISTS mykey
```

```
127.0.0.1:6666> scan 0 match "project:*
1) "17"
2) 1) "project:110"
   2) "project:109"
127.0.0.1:6666>
```

```
127.0.0.1:6666> scan 17 match *
1) "23"
2) 1) "run:id:734efc2df9314e33b913452176b76357"
   2) "run:project:110"
   3) "run:id:5e529f8d0c5546e3a2e0a79399437efe"
   4) "user"
   5) "run:id:d51123d53916478282bdbc7a15162d74"
   6) "run:id:59cf42018461423c89a6215d8001bdab"
   7) "project:46"
   8) "project:51"
   9) "run:id:00d54a4959844109880156233f5fe54b"
  10) "run:id:5b3d629744be4193bba30d68f39f4824"
127.0.0.1:6666>
```

```
127.0.0.1:6666> type run:id:8ceb7b0af5eb4836ae96e033d16ededeae  
hash
```

```
127.0.0.1:6666> hgetall run:id:8ceb7b0af5eb4836ae96e033d16ededeae  
1) "id"  
2) "8ceb7b0af5eb4836ae96e033d16ededeae"  
3) "p_id"  
4) "51"  
5) "parameters"  
6) "{\"begin\":-1.0,\"end\":-1.0}"  
7) "snapshottype"  
8) "brief"  
9) "status"  
10) "completed"  
11) "last_updated"  
12) "1713695310.318542"  
127.0.0.1:6666>
```

```
22) "run:id:59cf42018461423c89a6215d8001bdab"  
23) "run:id:892edad6c9ba47bcbbf2c6119fc7995a"  
24) "project:109"  
25) "run:id:a12a3080aeb1410b961ed28de323cb68"  
127.0.0.1:6666> type run:project:46  
list  
127.0.0.1:6666> lrange run:project:46 0 -1  
1) "c3e47ddd9b794f56a379f779088bb910"  
2) "a49ad162270a4fe497833c4281b9a6bf"  
3) "1e25b2dece464359b8b05c1060aabb98"  
127.0.0.1:6666> hgetall run:1e25b2dece464359b8b05c1060aabb98  
(empty array)  
127.0.0.1:6666> hgetall run:id:1e25b2dece464359b8b05c1060aabb98  
1) "id"  
2) "1e25b2dece464359b8b05c1060aabb98"  
3) "p_id"  
4) "46"  
5) "parameters"  
6) "{\"begin\":1711542601.0,\"end\":1713772955.72234}"  
7) "snapshottype"  
8) "brief"  
9) "status"  
10) "completed"  
11) "last_updated"  
12) "1713772967.819626"  
127.0.0.1:6666>
```

List

```
127.0.0.1:6666> sadd team natawut john  
(integer) 2  
127.0.0.1:6666> scard team  
(integer) 2  
127.0.0.1:6666> sadd team natawut john  
(integer) 0  
127.0.0.1:6666> scard team  
(integer) 2  
127.0.0.1:6666> smembers team  
1) "natawut"  
2) "john"  
127.0.0.1:6666>
```

ກຳລັງ ແລ້ວ Redis ກົມມະໂຫຼດ ມີ error
ກົມມະໂຫຼດ Redis

List

- List of strings, sorted by insertion order
- Can be used as list, queue, stack
 - LPUSH mylist abc # mylist contains "abc"
 - LPUSH mylist xyz # mylist contains "xyz", "abc"
 - RPUSH mylist 123 # mylist contains "xyz", "abc", "123"
- Use cases
 - Queue
 - Timelines

LPUSH → push מעתה
RPUSH → push תרשים (במאזע)

List Commands

BLMOVE

LMOVE

LSET

BLMPOP

LMPOP

LTRIM

BLPOP

LPOP

RPOP

BRPOP

LPOS

RPOPLPUSH

BRPOPLPUSH

LPUSH

RPUSH

LINDEX

LPUSHX

RPUSHNX

LINSERT

LRANGE

LLEN

LREM

~~Set~~

- Powerful data types for unordered non-duplicated keys
- Support many set operations e.g. intersection, union, etc.

```
SADD user_set natawut
SCARD user_set
SMEMBERS user_set
```
- Use cases
 - Set of user profiles
 - Set of inappropriate words for inappropriate content filtering

Set Commands

SADD

SISMEMBER

SSCAN

SCARD

SMEMBERS

SUNION

SDIFF

SMISMEMBER

SUNIONSTORE

SDIFFSTORE

SMOVE

SINTER

SPOP

SINTERCARD

SRANDMEMBER

SINTERSTORE

SREM

Sorted Set

ເຈັບຕາມ weight ໃນຮູບແບບກົດປົກກົດ

- Set of sorted items based on the score associated to each member

```
ZADD my_sortedset 5 data1
```

```
ZADD my_sortedset 1 data2 10 data3
```

```
ZRANGEBYSCORE my_sortedset 5. +inf WITHSCORES
```

- Use
 - Leader scoreboard
 - Priority queue

Sorted Set Commands

BZMPOP

ZDIFFSTORE

ZMSCORE

BZPOPMAX

ZINCRBY

ZPOPMAX

BZPOPMIN

ZINTER

ZPOPMIN

ZADD

ZINTERCARD

ZRANDMEMBER

ZCARD

ZINTERSTORE

ZRANGE

ZCOUNT

ZLEXCOUNT

ZRANGEBYLEX

ZDIFF

ZMPOP

ZRANGEBYSCORE

Hash

set 有多 data និង value នៃ string របស់ខ្លួន

- A container of unique fields and their values

HMSET profile:12345 user nnp id 12345 name "Natawut Nupairoj"
balance 10 key value

HGETALL profile:12345

HINCRBYprofile:12345 balance 5

- Use
 - User profile information
 - Post. Information

Hash Commands

HDEL

HLEN

HSTRLEN

HEXISTS

HMGET

HVALS

HGET

HMSET

HGETALL

HRANDFIELD

HINCRBY

HSCAN

HINCRBYFLOAT

HSET

HKEYS

HSETNX

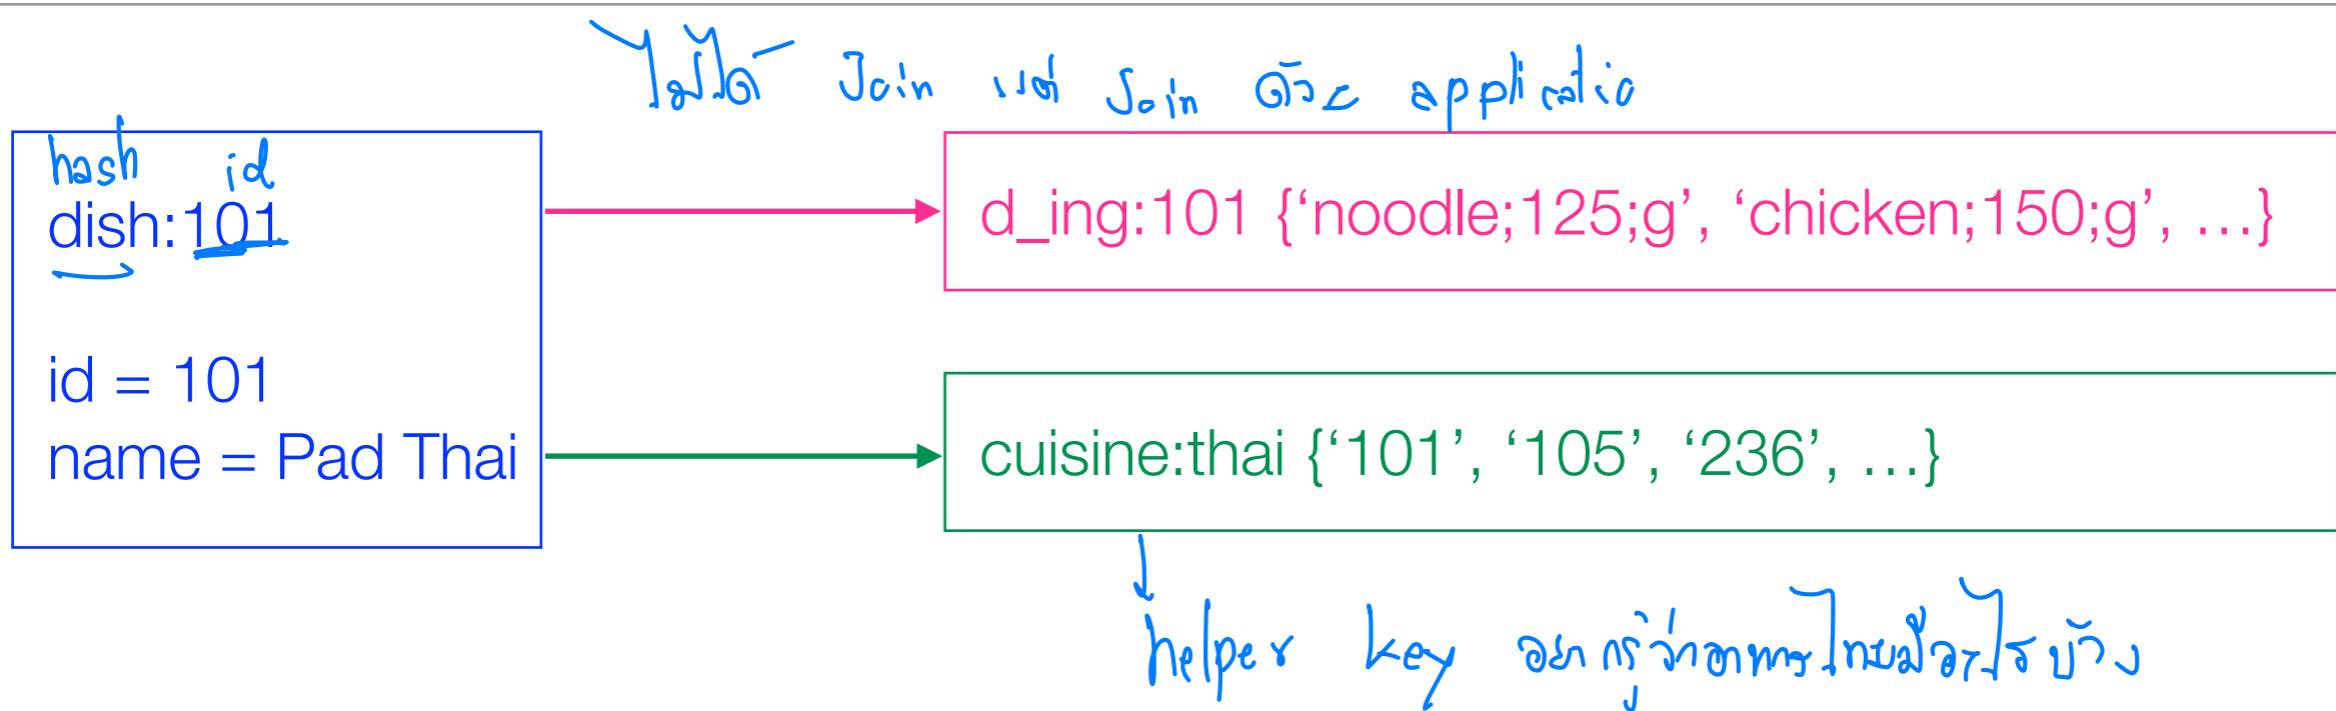
Example: Recipes in Redis

- Create data model for recipe
 - Data items
 - Main data
 - Dish - hash - key = dish:dish_id, value = dish_id, name, category, cuisine
 - Dish Ingredients - set - key = d_ing:dish_id, value = string(name; size; unit)
 - For query
 - Cuisine - set - key = cuisine:cuisine_value, value = string(dish_id)
 - Relationship
 - 1 dish can have multiple ingredients
 - Cuisine, as a helper data item, points back to dishes belong in the cuisine
- This is only one example of model design

take over คือการรับงาน

Example: Recipes in Redis

การหั่นชื่อ key สำหรับมาก เนื่องด้วย value ไม่ใช่



HMSET dish:101 id:101 name "Pad Thai" cuisine "thai" category "noodle"

SADD d_ing:101 "noodle;125;g"

SADD d_ing:101 "shrimp;150;g"

SADD d_ing:101 "brown sugar;3;tbsp"

SADD cuisine:thai "101"

จัดการข้อมูลร่วมกัน Redis ผ่าน API ของภาษา

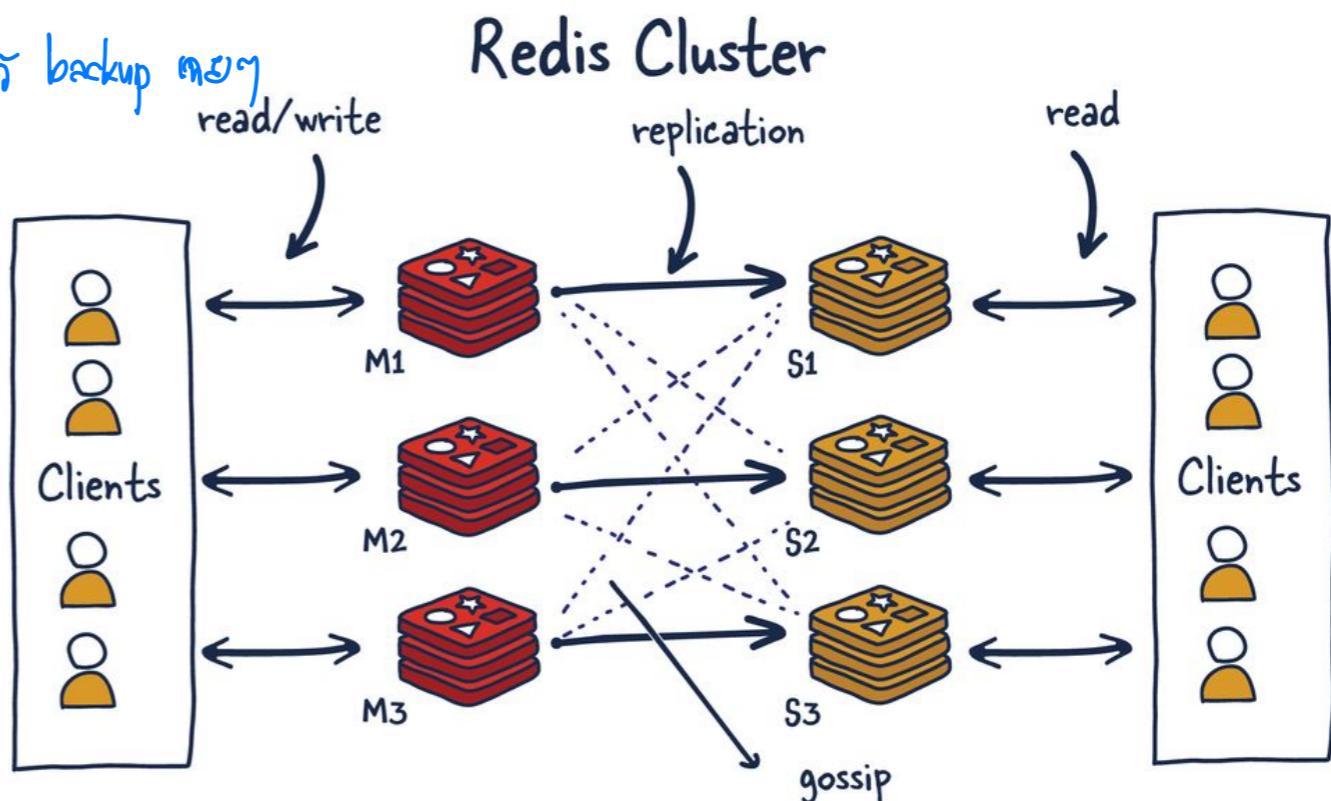
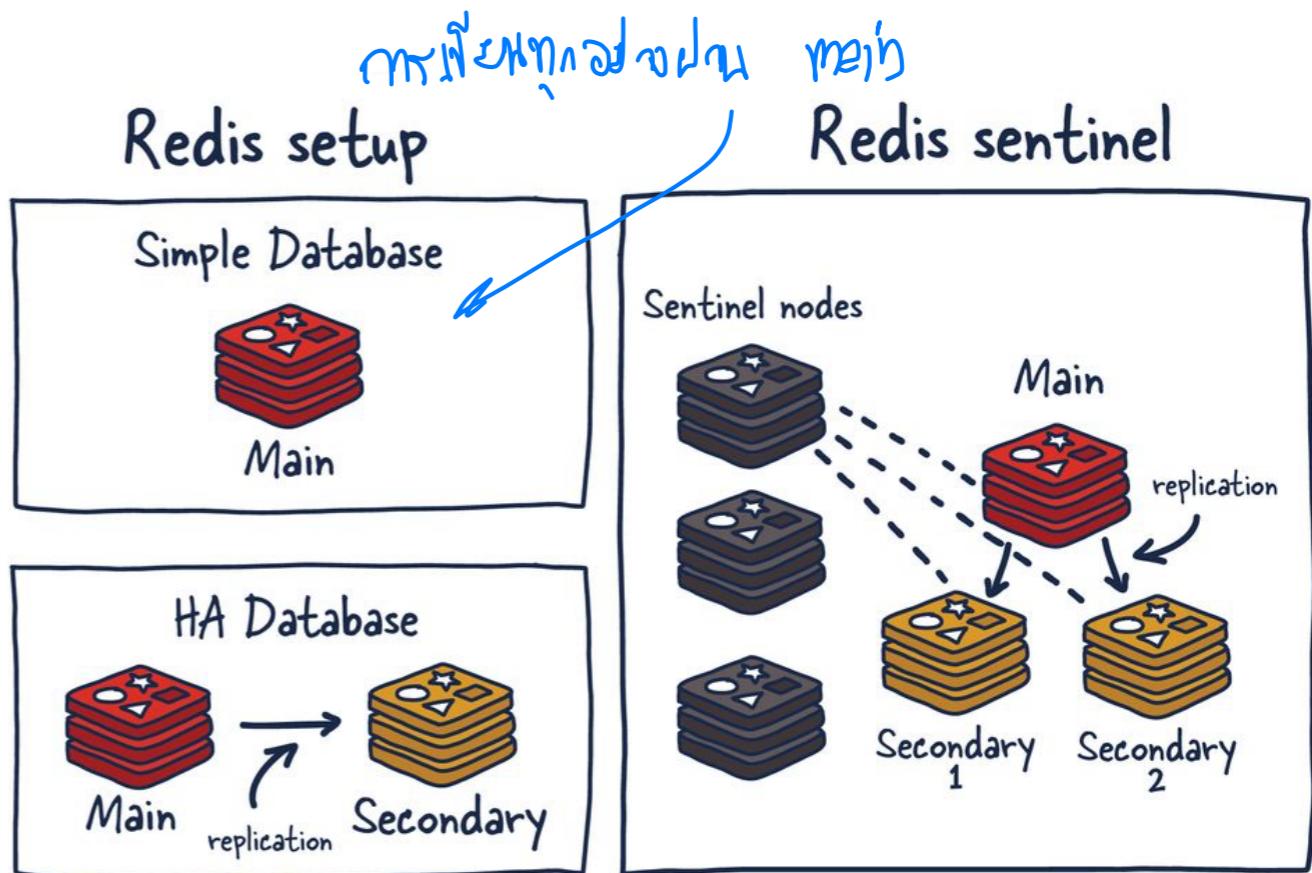
1. fail over

2. clustered

Redis Architecture

- Redis supports flexible architecture
 - Single instance
 - High Availability
 - Sentinel
 - Cluster

ຮັບຮູບທີ່ກຳນົດກາຍຕຸ້ນ ລາຍລະອຽດ



Redis High Availability

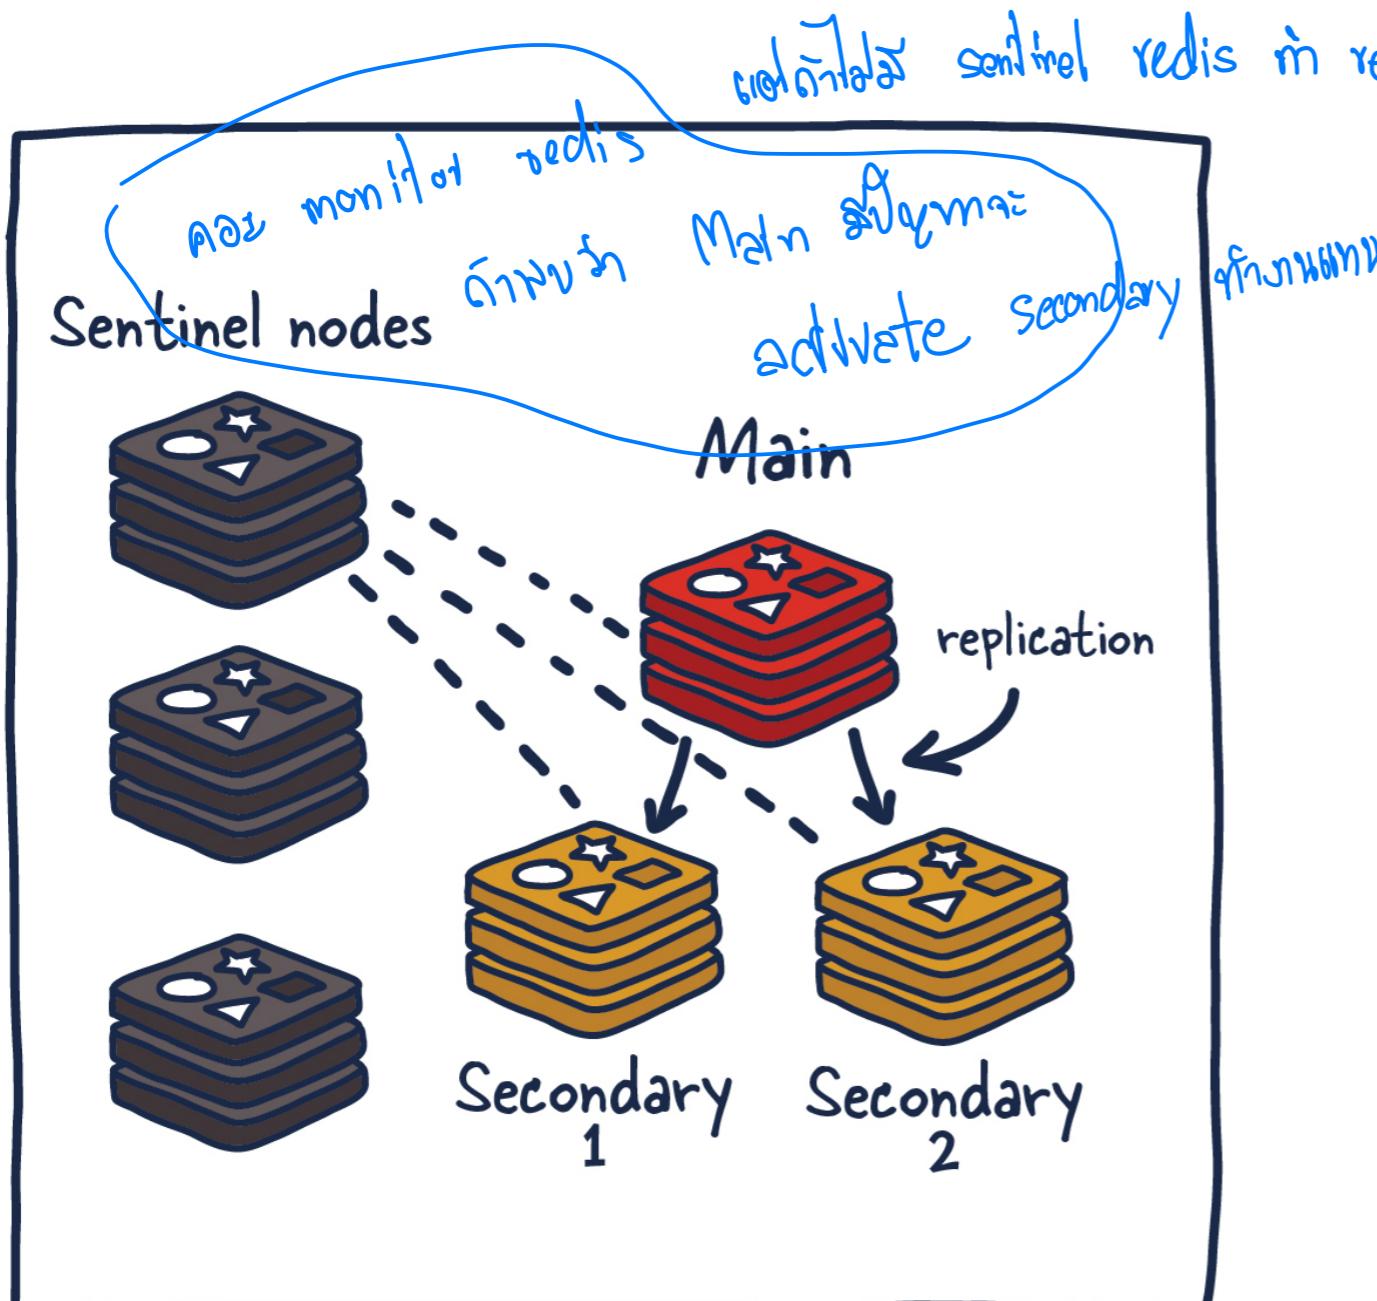
- When setup for HA, Redis consists of a master and one or more replicas and follows the following rules
 - Master performs all commands (read/write/etc.), replica performs only read commands
 - When a master and a replica instances are well-connected, the master keeps the replica updated by sending a stream of commands to the replica to replicate the effects on the dataset
 - When the link between the master and the replica breaks, for network issues or because a timeout is sensed in the master or the replica, the replica reconnects and attempts to proceed with a partial resynchronization
 - When a partial resynchronization is not possible, the replica will ask for a full resynchronization
 - Redis uses asynchronous replication, with asynchronous replica-to-master acknowledges of the amount of data processed

Redis High Availability

- To ensure data synchronization, all servers keep data version (replication id, offset); equal data versions mean both servers hold the same data
 - When a server becomes a master, it generates a new replication ID; replica server inherits replication ID from its master once connected
 - When master updates its data, it increments offset by one
- When replicas connect to masters, they use the PSYNC command to send their old master replication ID and the offsets they processed so far
 - If the differences of offsets are not too far behind, partial synchronization is performed
 - If too far or mismatched replication ID, a full resynchronization happens: in this case the replica will get a full copy of the dataset, from scratch

ສຳເນົາ ດ້ວຍ main ແລ້ວເປັນ replica ມີຄໍາໄວ້ takeover

Redis Sentinel



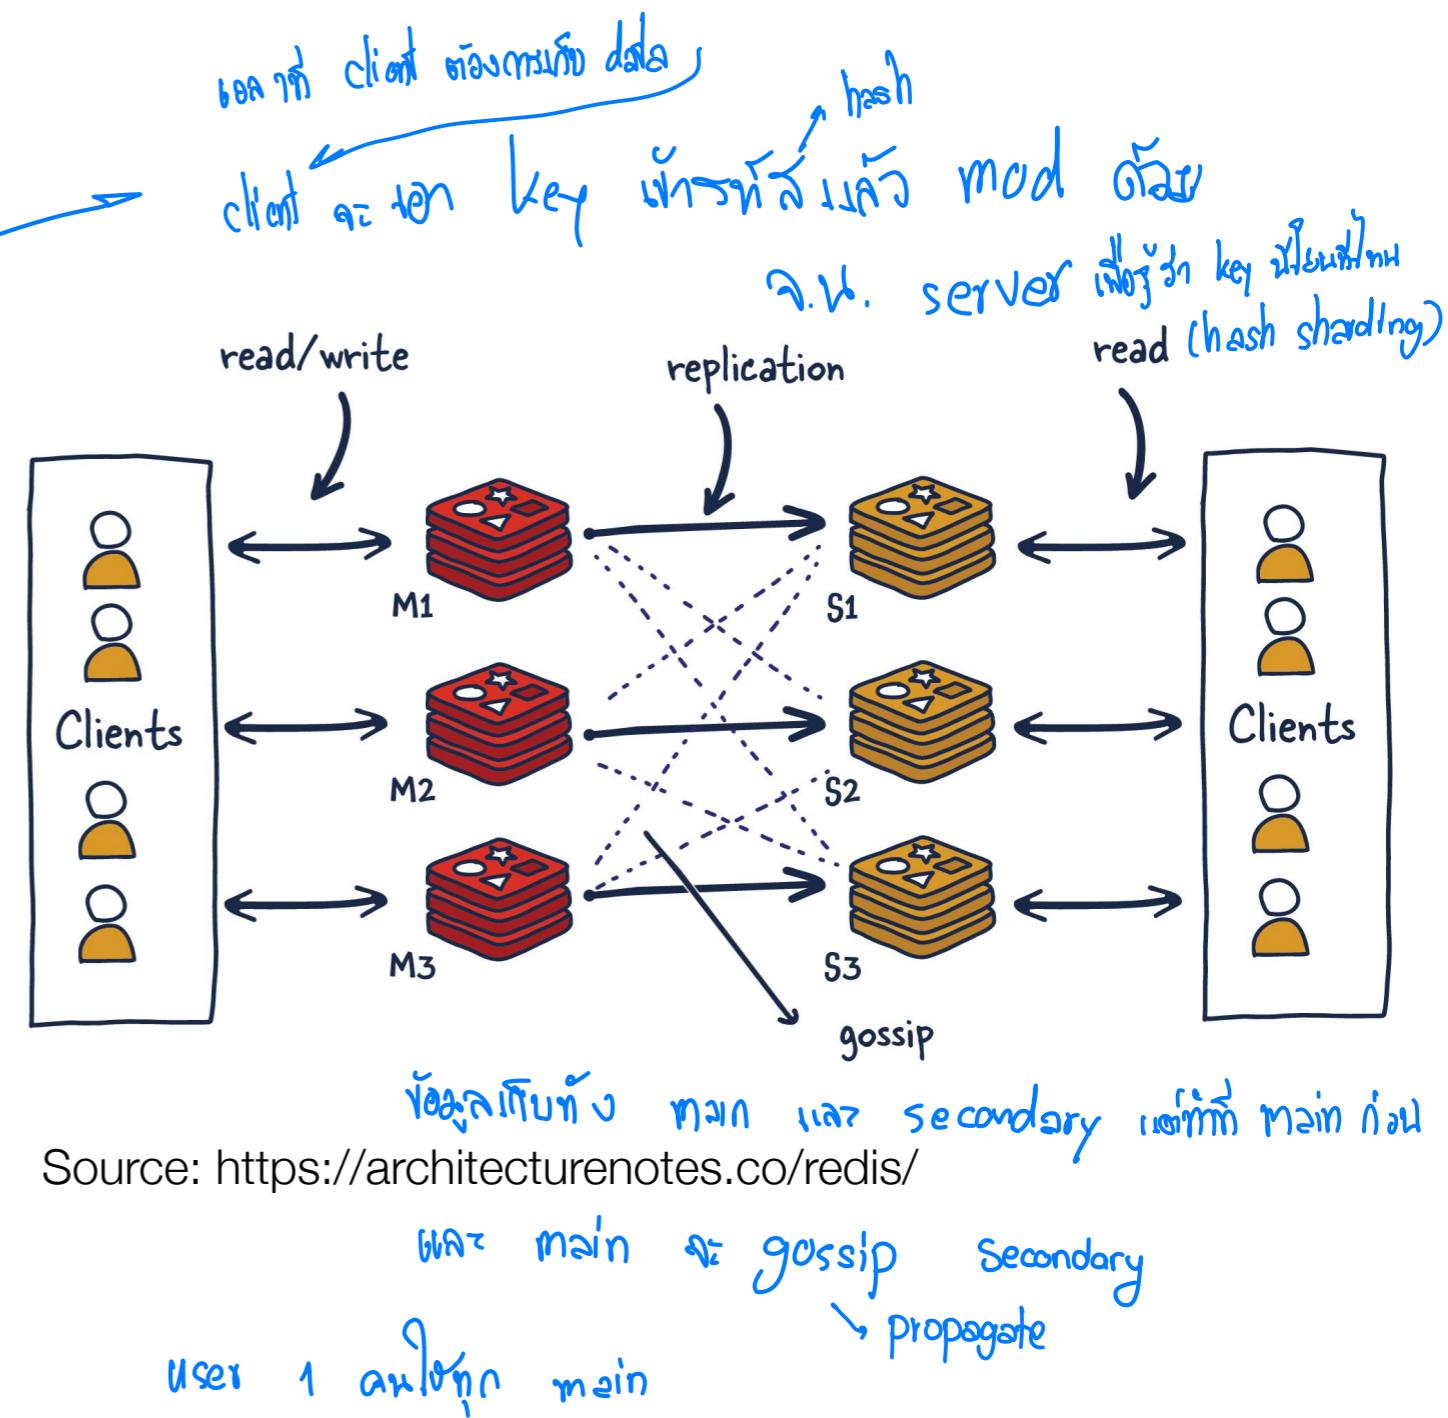
- A cluster of sentinel processes working together to coordinate states of Redis master (main) and replicas (secondaries)
- Focus on HA aspects with minimum number of nodes
- Support single master with multiple replicas with monitoring, notification, auto-failover, configuration providing
- Need at least three Sentinel instances for a robust deployment (quorum or agreement protocol)
- Can be deployed in Redis and client nodes

Source: <https://architecturenotes.co/redis/>

ផ្លូវ main នៃយុទ្ធណែលកំណត់ (scale out)

Redis Cluster

- Support algorithmic sharding for horizontal scaling
 - Incoming data is hashed into a hash slot using deterministic hash function and mod with number of nodes
 - Data with a given key always maps to the same shard *map ≈ hash slot*
 - Perform resharding when adding new nodes
- There are 16k hashslot; all data in the same hashslot always moves to the new node when resharding *មុនពីរដាក់ថ្មី នៅក្នុងថ្មី*
- Nodes communicate with one another using gossiping protocol to determine the entire cluster's health
- Redis Cluster does not guarantee strong consistency



Source: <https://architecturenotes.co/redis/>

References

- ScaleGrid, “Top Redis Use Cases by Core Data Structure Types”, <https://scalegrid.io/blog/top-redis-use-cases-by-core-data-structure-types/>
- Jerry An, “The most important Redis data structures you must understand”, <https://medium.com/analytics-vidhya/the-most-important-redis-data-structures-you-must-understand-2e95b5cf2bce>
- Brad Solomon, “How to use Redis with python”, <https://realpython.com/python-redis/>
- M. Yusuf, “Redis Explained”, <https://architecturenotes.co/redis/>